

In [1]:

```
from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force\_remount=True).

# Quora Question Pairs

## 1. Business Problem

### 1.1 Description </h2>

Quora is a place to gain and share knowledge—about anything. It's a platform to ask questions and connect with people who contribute unique insights and quality answers. This empowers people to learn from each other and to better understand the world.

Over 100 million people visit Quora every month, so it's no surprise that many people ask similarly worded questions. Multiple questions with the same intent can cause seekers to spend more time finding the best answer to their question, and make writers feel they need to answer multiple versions of the same question. Quora values canonical questions because they provide a better experience to active seekers and writers, and offer more value to both of these groups in the long term.

Credits: Kaggle

#### Problem Statement

- Identify which questions asked on Quora are duplicates of questions that have already been asked.
- This could be useful to instantly provide answers to questions that have already been answered.
- We are tasked with predicting whether a pair of questions are duplicates or not.

### 1.2 Sources/Useful Links</h2>

- Source : <https://www.kaggle.com/c/quora-question-pairs>

#### Useful Links

- Discussions : <https://www.kaggle.com/anokas/data-analysis-xgboost-starter-0-35460-lb/comments>
- Kaggle Winning Solution and other approaches:  
<https://www.dropbox.com/sh/93968nfnrzh8bp5/AACZdtsApc1QSTQc7X0H3QZ5a?dl=0>
- Blog 1 : <https://engineering.quora.com/Semantic-Question-Matching-with-Deep-Learning>
- Blog 2 : <https://towardsdatascience.com/identifying-duplicate-questions-on-quora-top-12-on-kaggle-4c1cf93f1c30>

### 1.3 Real world/Business Objectives and Constraints </h2>

1. The cost of a mis-classification can be very high.
2. You would want a probability of a pair of questions to be duplicates so that you can choose any threshold of choice.
3. No strict latency concerns.
4. Interpretability is partially important.

## 2. Machine Learning Problem

### 2.1 Data

#### 2.1.1 Data Overview

- Data will be in a file Train.csv
- Train.csv contains 5 columns : qid1, qid2, question1, question2, is\_duplicate
- Size of Train.csv - 60MB
- Number of rows in Train.csv = 404,290

#### 2.1.2 Example Data point

```
"id","qid1","qid2","question1","question2","is_duplicate"  
"0","1","2","What is the step by step guide to invest in share market in india?","What is  
the step by step guide to invest in share market?","0"  
"1","3","4","What is the story of Kohinoor (Koh-i-Noor) Diamond?","What would happen if the  
Indian government stole the Kohinoor (Koh-i-Noor) diamond back?","0"  
"7","15","16","How can I be a good geologist?","What should I do to be a great  
geologist?","1"  
"11","23","24","How do I read and find my YouTube comments?","How can I see all my Youtube  
comments?","1"
```

## 2.2 Mapping the real world problem to an ML problem

### 2.2.1 Type of Machine Learning Problem

It is a binary classification problem, for a given pair of questions we need to predict if they are duplicate or not.

### 2.2.2 Performance Metric

Source: <https://www.kaggle.com/c/quora-question-pairs#evaluation>

Metric(s):

- log-loss : <https://www.kaggle.com/wiki/LogarithmicLoss>
- Binary Confusion Matrix

## 2.3 Train and Test Construction

We build train and test by randomly splitting in the ratio of 70:30 or 80:20 whatever we choose as we have sufficient points to work with.

## 3. Exploratory Data Analysis

In [2]:

```
file_path = "/content/drive/MyDrive/Applied AI Assignments/LN Assignments/Case Study 1 Quora question pair/"  
  
#"/content/drive/MyDrive/Applied AI Assignments/case_study_1_Quora/"
```

In [3]:

```
pip install distance
```

Requirement already satisfied: distance in /usr/local/lib/python3.7/dist-packages (0.1.3)

In [4]:

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from subprocess import check_output
%matplotlib inline
import plotly.offline as py
py.init_notebook_mode(connected=True)
import plotly.graph_objs as go
import plotly.tools as tls
import os
import gc

import re
from nltk.corpus import stopwords
import distance
from nltk.stem import PorterStemmer
from bs4 import BeautifulSoup
```

## 3.1 Reading data and basic stats

In [5]:

```
df = pd.read_csv(file_path + "original_train.csv")

print("Number of data points:", df.shape[0])
```

Number of data points: 404290

In [6]:

```
df = df.sample(n = 10000)
```

In [7]:

```
df.to_csv(file_path + "train.csv", index=False)
```

In [8]:

```
df['id'].shape
```

Out[8]:

(10000,)

In [9]:

```
df.head()
```

Out[9]:

	id	qid1	qid2	question1	question2	is_duplicate
206552	206552	69616	310006	Where can I find free TV shows online without ...	Where can I use a credit card without the CVV ...	0
108228	108228	58375	177877	Who is going to Delhi this year for IAS prepar...	Which is the best month to go to Delhi for IAS...	0
13096	13096	25172	25173	Why is the name of Hashem (G-D) not mentioned ...	Why isn't God mentioned in the Book of Esther?	1

id	qid1	qid2	question1	question2	is_duplicate
173278	173278	68426	How do I overcome from depression without getting...	How is depression cured with therapist?	1
174488	174488	268903	How is life as a technical support engineer at...	Will Microsoft GTSC help me build my career in...	0

In [10]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 10000 entries, 206552 to 165297
Data columns (total 6 columns):
#   Column          Non-Null Count  Dtype
---  -
0   id              10000 non-null   int64
1   qid1            10000 non-null   int64
2   qid2            10000 non-null   int64
3   question1       10000 non-null   object
4   question2       10000 non-null   object
5   is_duplicate    10000 non-null   int64
dtypes: int64(4), object(2)
memory usage: 546.9+ KB
```

We are given a minimal number of data fields here, consisting of:

- id: Looks like a simple rowID
- qid{1, 2}: The unique ID of each question in the pair
- question{1, 2}: The actual textual contents of the questions.
- is\_duplicate: The label that we are trying to predict - whether the two questions are duplicates of each other.

### 3.2.1 Distribution of data points among output classes


- Number of duplicate(similar) and non-duplicate(non similar) questions

In [11]:

```
df.groupby("is_duplicate")["id"].count().plot.bar()
```

Out[11]:

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f069682d750>
```



In [12]:

```
print('~> Total number of question pairs for training:\n {}'.format(len(df)))
```

```
~> Total number of question pairs for training:
10000
```

In [13]:

```
print('~> Question pairs are not Similar (is_duplicate = 0):\n    {}'.format(100 -
round(df['is_duplicate'].mean()*100, 2)))
print('\n~> Question pairs are Similar (is_duplicate = 1):\n    {}'.format(round(df['is_duplicate']
].mean()*100, 2)))
```

```
~> Question pairs are not Similar (is_duplicate = 0):
    62.81%
```

```
~> Question pairs are Similar (is_duplicate = 1):
    37.19%
```

### 3.2.2 Number of unique questions </h3>

In [14]:

```
qids = pd.Series(df['qid1'].tolist() + df['qid2'].tolist())
unique_qs = len(np.unique(qids))
qs_morethan_onetime = np.sum(qids.value_counts() > 1)
print ('Total number of Unique Questions are: {}'.format(unique_qs))
#print len(np.unique(qids))

print ('Number of unique questions that appear more than one time: {}
({})%\n'.format(qs_morethan_onetime,qs_morethan_onetime/unique_qs*100))

print ('Max number of times a single question is repeated: {}'.format(max(qids.value_counts()))))

q_vals=qids.value_counts()

q_vals=q_vals.values
```

Total number of Unique Questions are: 19347

Number of unique questions that appear more than one time: 575 (2.9720370083217036%)

Max number of times a single question is repeated: 5


In [15]:


```
x = ["unique_questions" , "Repeated Questions"]
y = [unique_qs , qs_morethan_onetime]

plt.figure(figsize=(10, 6))
plt.title ("Plot representing unique and repeated questions ")
sns.barplot(x,y)
plt.show()
```

/usr/local/lib/python3.7/dist-packages/seaborn/\_decorators.py:43: FutureWarning:

Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.





### 3.2.3 Checking for Duplicates

In [16]:

```
#checking whether there are any repeated pair of questions

pair_duplicates =
df[['qid1', 'qid2', 'is_duplicate']].groupby(['qid1', 'qid2']).count().reset_index()

print ("Number of duplicate questions", (pair_duplicates).shape[0] - df.shape[0])
```

Number of duplicate questions 0

### 3.2.4 Number of occurrences of each question

In [17]:

```
plt.figure(figsize=(20, 10))

plt.hist(qids.value_counts(), bins=160)

plt.yscale('log', nonposy='clip')


plt.title('Log-Histogram of question appearance counts')

plt.xlabel('Number of occurrences of question')

plt.ylabel('Number of questions')

print ('Maximum number of times a single question is repeated: {}'.format(max(qids.value_counts(
))))
```

Maximum number of times a single question is repeated: 5



### 3.2.5 Checking for NULL values

### 3.2.3 Checking for NULL values </h3>

In [18]:

```
#Checking whether there are any rows with null values
nan_rows = df[df.isnull().any(1)]
print (nan_rows)
```

Empty DataFrame

Columns: [id, qid1, qid2, question1, question2, is\_duplicate]

Index: []

- There are two rows with null values in question2

In [19]:

```
# Filling the null values with ' '
df = df.fillna(' ')
nan_rows = df[df.isnull().any(1)]
print (nan_rows)
```

Empty DataFrame

Columns: [id, qid1, qid2, question1, question2, is\_duplicate]

Index: []

## 3.3 Basic Feature Extraction (before cleaning) </h2>

Let us now construct a few features like:

- **freq\_qid1** = Frequency of qid1's
- **freq\_qid2** = Frequency of qid2's
- **q1len** = Length of q1
- **q2len** = Length of q2
- **q1\_n\_words** = Number of words in Question 1
- **q2\_n\_words** = Number of words in Question 2
- **word\_Common** = (Number of common unique words in Question 1 and Question 2)
- **word\_Total** = (Total num of words in Question 1 + Total num of words in Question 2)
- **word\_share** = (word\_common)/(word\_Total)
- **freq\_q1+freq\_q2** = sum total of frequency of qid1 and qid2
- **freq\_q1-freq\_q2** = absolute difference of frequency of qid1 and qid2

In [20]:

```
if os.path.isfile(file_path+'df_fe_without_preprocessing_train.csv'):
    df = pd.read_csv(file_path+"df_fe_without_preprocessing_train.csv",encoding='latin-1')
else:
    df['freq_qid1'] = df.groupby('qid1')['qid1'].transform('count')
    df['freq_qid2'] = df.groupby('qid2')['qid2'].transform('count')
    df['q1len'] = df['question1'].str.len()
    df['q2len'] = df['question2'].str.len()
    df['q1_n_words'] = df['question1'].apply(lambda row: len(row.split(" ")))
    df['q2_n_words'] = df['question2'].apply(lambda row: len(row.split(" ")))

    def normalized_word_Common(row):
        w1 = set(map(lambda word: word.lower().strip(), row['question1'].split(" ")))
        w2 = set(map(lambda word: word.lower().strip(), row['question2'].split(" ")))
        return 1.0 * len(w1 & w2)
    df['word_Common'] = df.apply(normalized_word_Common, axis=1)

    def normalized_word_Total(row):
        w1 = set(map(lambda word: word.lower().strip(), row['question1'].split(" ")))
        w2 = set(map(lambda word: word.lower().strip(), row['question2'].split(" ")))
        return 1.0 * (len(w1) + len(w2))
    df['word_Total'] = df.apply(normalized_word_Total, axis=1)

    def normalized_word_share(row):
        w1 = set(map(lambda word: word.lower().strip(), row['question1'].split(" ")))
        w2 = set(map(lambda word: word.lower().strip(), row['question2'].split(" ")))
        return 1.0 * len(w1 & w2) / (len(w1) + len(w2))
    df['word_share'] = df.apply(normalized_word_share, axis=1)
```

```
df['word_share'] = df.apply(normalized_word_share, axis=1)

df['freq_q1+q2'] = df['freq_qid1']+df['freq_qid2']
df['freq_q1-q2'] = abs(df['freq_qid1']-df['freq_qid2'])

df.to_csv(file_path+"df_fe_without_preprocessing_train.csv", index=False)

df.head()
```

Out[20]:

	id	qid1	qid2	question1	question2	is_duplicate	freq_qid1	freq_qid2	q1len	q2len	q1_n_words	q2_n
<b>206552</b>	206552	69616	310006	Where can I find free TV shows online without ...	Where can I use a credit card without the CVV ...	0	1	1	74	54	14	11
<b>108228</b>	108228	58375	177877	Who is going to Delhi this year for IAS prepar...	Which is the best month to go to Delhi for IAS...	0	1	1	52	59	10	12
<b>13096</b>	13096	25172	25173	Why is the name of Hashem (G-D) not mentioned ...	Why isn't God mentioned in the Book of Esther?	1	1	1	72	46	14	9
<b>173278</b>	173278	68426	30173	How do I overcome from depression without gett...	How is depression cured without a therapist?	1	1	1	67	44	11	7
<b>174488</b>	174488	268903	268904	How is life as a technical support engineer at...	Will Microsoft GTSC help me build my career in...	0	1	1	62	58	11	10

In [21]:

```
df['id'].shape
```

Out[21]:

```
(10000,)
```

### 3.3.1 Analysis of some of the extracted features

- Here are some questions have only one single words.

In [22]:

```
print ("Minimum length of the questions in question1 : " , min(df['q1_n_words']))

print ("Minimum length of the questions in question2 : " , min(df['q2_n_words']))

print ("Number of Questions with minimum length [question1] :", df[df['q1_n_words']== 1].shape[0])
print ("Number of Questions with minimum length [question2] :", df[df['q2_n_words']== 1].shape[0])
```



```
Minimum length of the questions in question1 : 1
Minimum length of the questions in question2 : 1
Number of Questions with minimum length [question1] : 3
Number of Questions with minimum length [question2] : 1
```

### 3.3.1.1 Feature: word\_share </h4>

In [23]:

```
plt.figure(figsize=(12, 8))

plt.subplot(1,2,1)
sns.violinplot(x = 'is_duplicate', y = 'word_share', data = df[0:])


plt.subplot(1,2,2)
sns.distplot(df[df['is_duplicate'] == 1.0]['word_share'][0:], label = "1", color = 'red')
sns.distplot(df[df['is_duplicate'] == 0.0]['word_share'][0:], label = "0", color = 'blue' )
plt.show()
```

/usr/local/lib/python3.7/dist-packages/seaborn/distributions.py:2557: FutureWarning:

`distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

/usr/local/lib/python3.7/dist-packages/seaborn/distributions.py:2557: FutureWarning:

`distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).



- The distributions for normalized word\_share have some overlap on the far right-hand side, i.e., there are quite a lot of questions with high word similarity
- The average word share and Common no. of words of qid1 and qid2 is more when they are duplicate(Similar)

### 3.3.1.2 Feature: word\_Common </h4>

In [24]:

```
plt.figure(figsize=(12, 8))

plt.subplot(1,2,1)
sns.violinplot(x = 'is_duplicate', y = 'word_Common', data = df[0:])


plt.subplot(1,2,2)
sns.distplot(df[df['is_duplicate'] == 1.0]['word_Common'][0:] , label = "1", color = 'red')
sns.distplot(df[df['is_duplicate'] == 0.0]['word_Common'][0:] , label = "0" , color = 'blue' )
plt.show()
```

/usr/local/lib/python3.7/dist-packages/seaborn/distributions.py:2557: FutureWarning:

`distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

/usr/local/lib/python3.7/dist-packages/seaborn/distributions.py:2557: FutureWarning:

`distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).



The distributions of the word\_Common feature in similar and non-similar questions are highly overlapping

## 4.EDA: Advanced Feature Extraction.

In [25]:

```
pip install fuzzywuzzy
```

Requirement already satisfied: fuzzywuzzy in /usr/local/lib/python3.7/dist-packages (0.18.0)

In [26]:

```
import warnings
warnings.filterwarnings("ignore")
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from subprocess import check_output
```

```
%matplotlib inline
import plotly.offline as py
py.init_notebook_mode(connected=True)
import plotly.graph_objs as go
import plotly.tools as tls
import os
import gc

import re
from nltk.corpus import stopwords
import distance
from nltk.stem import PorterStemmer
from bs4 import BeautifulSoup
import re
from nltk.corpus import stopwords
# This package is used for finding longest common subsequence between two strings
# you can write your own dp code for this
import distance
from nltk.stem import PorterStemmer
from bs4 import BeautifulSoup
from fuzzywuzzy import fuzz
from sklearn.manifold import TSNE
# Import the Required lib packages for WORD-Cloud generation
# https://stackoverflow.com/questions/45625434/how-to-install-wordcloud-in-python3-6
from wordcloud import WordCloud, STOPWORDS
from os import path
from PIL import Image
```

In [27]:

```
df.head(2)
```

Out[27]:

	id	qid1	qid2	question1	question2	is_duplicate	freq_qid1	freq_qid2	q1len	q2len	q1_n_words	q2_n_v
206552	206552	69616	310006	Where can I find free TV shows online without ...	Where can I use a credit card without the CVV ...	0	1	1	74	54	14	11
108228	108228	58375	177877	Who is going to Delhi this year for IAS prepar...	Which is the best month to go to Delhi for IAS...	0	1	1	52	59	10	12

In [28]:

```
df['id'].shape
```

Out[28]:

```
(10000,)
```

## 4.1 Preprocessing of Text

- Preprocessing:
  - Removing html tags
  - Removing Punctuations
  - Performing stemming
  - Removing Stopwords

- Expanding contractions etc.

In [29]:

```
import nltk
nltk.download('stopwords')
```

```
[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data]   Package stopwords is already up-to-date!
```

Out[29]:

True

In [30]:

```
# To get the results in 4 decemal points
SAFE_DIV = 0.0001

STOP_WORDS = stopwords.words("english")

def preprocess(x):
    x = str(x).lower()
    x = x.replace(",000,000", "m").replace(",000", "k").replace("'", "").replace('"', "")\
        .replace("won't", "will not").replace("cannot", "can not").replace("can't", "can not")\
        .replace("n't", " not").replace("what's", "what is").replace("it's", "it s")\
        .replace("'ve", " have").replace("i'm", "i am").replace("'re", " are")\
        .replace("he's", "he is").replace("she's", "she is").replace("'s", " own"
    \
        .replace("%", " percent ").replace("₹", " rupee ").replace("$", " dollar
    )\
        .replace("€", " euro ").replace("'ll", " will")
    x = re.sub(r"([0-9]+)000000", r"\1m", x)
    x = re.sub(r"([0-9]+)000", r"\1k", x)

    porter = PorterStemmer()
    pattern = re.compile('\W')

    if type(x) == type(''):
        x = re.sub(pattern, ' ', x)

    if type(x) == type(''):
        x = porter.stem(x)
        example1 = BeautifulSoup(x)
        x = example1.get_text()

    return x
```

- Function to Compute and get the features : With 2 parameters of Question 1 and Question 2

## 4.2 Advanced Feature Extraction (NLP and Fuzzy Features)

Definition:

- **Token**: You get a token by splitting sentence a space
- **Stop\_Word** : stop words as per NLTK.
- **Word** : A token that is not a stop\_word

Features:

- **cwc\_min** : Ratio of common\_word\_count to min length of word count of Q1 and Q2  

$$\text{cwc\_min} = \text{common\_word\_count} / (\min(\text{len}(q1\_words), \text{len}(q2\_words)))$$
- **cwc\_max** : Ratio of common\_word\_count to max length of word count of Q1 and Q2  

$$\text{cwc\_max} = \text{common\_word\_count} / (\max(\text{len}(q1\_words), \text{len}(q2\_words)))$$
- **csc\_min** : Ratio of common\_stop\_count to min length of stop count of Q1 and Q2

- $csc\_min = common\_stop\_count / (\min(len(q1\_stops), len(q2\_stops)))$
- **csc\_max** : Ratio of common\_stop\_count to max length of stop count of Q1 and Q2  
 $csc\_max = common\_stop\_count / (\max(len(q1\_stops), len(q2\_stops)))$
  - **ctc\_min** : Ratio of common\_token\_count to min length of token count of Q1 and Q2  
 $ctc\_min = common\_token\_count / (\min(len(q1\_tokens), len(q2\_tokens)))$
  - **ctc\_max** : Ratio of common\_token\_count to max length of token count of Q1 and Q2  
 $ctc\_max = common\_token\_count / (\max(len(q1\_tokens), len(q2\_tokens)))$
  - **last\_word\_eq** : Check if First word of both questions is equal or not  
 $last\_word\_eq = int(q1\_tokens[-1] == q2\_tokens[-1])$
  - **first\_word\_eq** : Check if First word of both questions is equal or not  
 $first\_word\_eq = int(q1\_tokens[0] == q2\_tokens[0])$
  - **abs\_len\_diff** : Abs. length difference  
 $abs\_len\_diff = abs(len(q1\_tokens) - len(q2\_tokens))$
  - **mean\_len** : Average Token Length of both Questions  
 $mean\_len = (len(q1\_tokens) + len(q2\_tokens))/2$
  - **fuzz\_ratio** : <https://github.com/seatgeek/fuzzywuzzy#usage> <http://chairnerd.seatgeek.com/fuzzywuzzy-fuzzy-string-matching-in-python/>
  - **fuzz\_partial\_ratio** : <https://github.com/seatgeek/fuzzywuzzy#usage> <http://chairnerd.seatgeek.com/fuzzywuzzy-fuzzy-string-matching-in-python/>
  - **token\_sort\_ratio** : <https://github.com/seatgeek/fuzzywuzzy#usage> <http://chairnerd.seatgeek.com/fuzzywuzzy-fuzzy-string-matching-in-python/>
  - **token\_set\_ratio** : <https://github.com/seatgeek/fuzzywuzzy#usage> <http://chairnerd.seatgeek.com/fuzzywuzzy-fuzzy-string-matching-in-python/>
  - **longest\_substr\_ratio** : Ratio of length longest common substring to min length of token count of Q1 and Q2  
 $longest\_substr\_ratio = len(longest\ common\ substring) / (\min(len(q1\_tokens), len(q2\_tokens)))$

In [31]:

```
def get_token_features(q1, q2):
    token_features = [0.0]*10

    # Converting the Sentence into Tokens:
    q1_tokens = q1.split()
    q2_tokens = q2.split()

    if len(q1_tokens) == 0 or len(q2_tokens) == 0:
        return token_features

    # Get the non-stopwords in Questions
    q1_words = set([word for word in q1_tokens if word not in STOP_WORDS])
    q2_words = set([word for word in q2_tokens if word not in STOP_WORDS])

    #Get the stopwords in Questions
    q1_stops = set([word for word in q1_tokens if word in STOP_WORDS])
    q2_stops = set([word for word in q2_tokens if word in STOP_WORDS])

    # Get the common non-stopwords from Question pair
    common_word_count = len(q1_words.intersection(q2_words))

    # Get the common stopwords from Question pair
    common_stop_count = len(q1_stops.intersection(q2_stops))

    # Get the common Tokens from Question pair
    common_token_count = len(set(q1_tokens).intersection(set(q2_tokens)))

    token_features[0] = common_word_count / (min(len(q1_words), len(q2_words)) + SAFE_DIV)
    token_features[1] = common_word_count / (max(len(q1_words), len(q2_words)) + SAFE_DIV)
    token_features[2] = common_stop_count / (min(len(q1_stops), len(q2_stops)) + SAFE_DIV)
    token_features[3] = common_stop_count / (max(len(q1_stops), len(q2_stops)) + SAFE_DIV)
    token_features[4] = common_token_count / (min(len(q1_tokens), len(q2_tokens)) + SAFE_DIV)
    token_features[5] = common_token_count / (max(len(q1_tokens), len(q2_tokens)) + SAFE_DIV)

    # Last word of both question is same or not
    token_features[6] = int(q1_tokens[-1] == q2_tokens[-1])

    # First word of both question is same or not
    token_features[7] = int(q1_tokens[0] == q2_tokens[0])

    token_features[8] = abs(len(q1_tokens) - len(q2_tokens))
```

```

#Average Token Length of both Questions
token_features[9] = (len(q1_tokens) + len(q2_tokens))/2
return token_features

# get the Longest Common sub string

def get_longest_substr_ratio(a, b):
    strs = list(distance.lcs substrings(a, b))
    if len(strs) == 0:
        return 0
    else:
        return len(strs[0]) / (min(len(a), len(b)) + 1)

def extract_features(df):
    # preprocessing each question
    df["question1"] = df["question1"].fillna("").apply(preprocess)
    df["question2"] = df["question2"].fillna("").apply(preprocess)

    print("token features...")

    # Merging Features with dataset

    token_features = df.apply(lambda x: get_token_features(x["question1"], x["question2"]), axis=1)
    # print(token_features)

    df["cwc_min"] = list(map(lambda x: x[0], token_features))
    df["cwc_max"] = list(map(lambda x: x[1], token_features))
    df["csc_min"] = list(map(lambda x: x[2], token_features))
    df["csc_max"] = list(map(lambda x: x[3], token_features))
    df["ctc_min"] = list(map(lambda x: x[4], token_features))
    df["ctc_max"] = list(map(lambda x: x[5], token_features))
    df["last_word_eq"] = list(map(lambda x: x[6], token_features))
    df["first_word_eq"] = list(map(lambda x: x[7], token_features))
    df["abs_len_diff"] = list(map(lambda x: x[8], token_features))
    df["mean_len"] = list(map(lambda x: x[9], token_features))

    #Computing Fuzzy Features and Merging with Dataset

    # do read this blog: http://chairnerd.seatgeek.com/fuzzywuzzy-fuzzy-string-matching-in-python/
    # https://stackoverflow.com/questions/31806695/when-to-use-which-fuzz-function-to-compare-2-strings
    # https://github.com/seatgeek/fuzzywuzzy
    print("fuzzy features..")

    df["token_set_ratio"] = df.apply(lambda x: fuzz.token_set_ratio(x["question1"],
x["question2"]), axis=1)
    # The token sort approach involves tokenizing the string in question, sorting the tokens alpha
betically, and
    # then joining them back into a string We then compare the transformed strings with a simple r
atio().
    df["token_sort_ratio"] = df.apply(lambda x: fuzz.token_sort_ratio(x["question1"],
x["question2"]), axis=1)
    df["fuzz_ratio"] = df.apply(lambda x: fuzz.QRatio(x["question1"], x["question2"]), a
is=1)
    df["fuzz_partial_ratio"] = df.apply(lambda x: fuzz.partial_ratio(x["question1"],
x["question2"]), axis=1)
    df["longest_substr_ratio"] = df.apply(lambda x: get_longest_substr_ratio(x["question1"], x["qu
estion2"]), axis=1)
    return df

```

In [32]:

```

# if os.path.isfile(file_path + "nlp_features_train.csv"): # file_path + "nlp_features_train.csv"
#     df = pd.read_csv(file_path + "nlp_features_train.csv", encoding='latin-1')
#     df.fillna('')
# else:
print("Extracting features for train:")
# df = pd.read_csv( file_path + "train.csv")
df = extract_features(df)
df.to_csv(file_path+"nlp_features_train.csv", index=False)

```

Extracting features for train:  
token features...  
fuzzy features..

In [33]:

```
df.columns
```

Out[33]:

```
Index(['id', 'qid1', 'qid2', 'question1', 'question2', 'is_duplicate',
      'freq_qid1', 'freq_qid2', 'q1len', 'q2len', 'q1_n_words', 'q2_n_words',
      'word_Common', 'word_Total', 'word_share', 'freq_q1+q2', 'freq_q1-q2',
      'cwc_min', 'cwc_max', 'csc_min', 'csc_max', 'ctc_min', 'ctc_max',
      'last_word_eq', 'first_word_eq', 'abs_len_diff', 'mean_len',
      'token_set_ratio', 'token_sort_ratio', 'fuzz_ratio',
      'fuzz_partial_ratio', 'longest_substr_ratio'],
      dtype='object')
```

In [34]:

```
df.head(2)
```

Out[34]:

	id	qid1	qid2	question1	question2	is_duplicate	freq_qid1	freq_qid2	q1len	q2len	q1_n_words	q2_n_v
206552	206552	69616	310006	where can i find free tv shows online without ...	where can i use a credit card without the cvv ...	0	1	1	74	54	14	11
108228	108228	58375	177877	who is going to delhi this year for ias prepar...	which is the best month to go to delhi for ias...	0	1	1	52	59	10	12

In [35]:

```
df['id'].shape
```

Out[35]:

```
(10000,)
```

## 4.2.1 Analysis of extracted features

### 4.2.1.1 Plotting Word clouds

- Creating Word Cloud of Duplicates and Non-Duplicates Question pairs
- We can observe the most frequent occurring words

In [36]:

```
df_duplicate = df[df['is_duplicate'] == 1]
dfp_nonduplicate = df[df['is_duplicate'] == 0]

# Converting 2d array of q1 and q2 and flatten the array: like {{1,2},{3,4}} to {1,2,3,4}
p = np.dstack([df_duplicate["question1"], df_duplicate["question2"]]).flatten()
n = np.dstack([dfp_nonduplicate["question1"], dfp_nonduplicate["question2"]]).flatten()

print ("Number of data points in class 1 (duplicate pairs) :",len(p))
print ("Number of data points in class 0 (non duplicate pairs) :",len(n))

#Saving the np array into a text file
np.savetxt(file path + 'train p.txt', p, delimiter=' ', fmt='%s')
```

```
Number of data points in class 1 (duplicate pairs) : 7438
Number of data points in class 0 (non duplicate pairs) : 12562
```

```
# reading the text files and removing the Stop Words:
d = path.dirname(file_path)

textp_w = open(path.join(d, 'train_p.txt')).read()
textn_w = open(path.join(d, 'train_n.txt')).read()
stopwords = set(STOPWORDS)
stopwords.add("said")
stopwords.add("br")
stopwords.add(" ")
stopwords.remove("not")

stopwords.remove("no")
#stopwords.remove("good")
#stopwords.remove("love")
stopwords.remove("like")
#stopwords.remove("best")
#stopwords.remove("!")

print ("Total number of words in duplicate pair questions :",len(textp_w))
print ("Total number of words in non duplicate pair questions :",len(textn_w))
```

**Word Clouds generated from duplicate pair question's text**

```

wc = WordCloud(background_color="white", max_words=len(textp_w), stopwords=stopwords)
wc.generate(textp_w)
print("Word Cloud for Duplicate Question pairs")
plt.imshow(wc, interpolation='bilinear')
plt.axis("off")
plt.show()

```

```
wc = WordCloud(background_color="white", max_words=len(textn_w), stopwords=stopwords)
# generate word cloud
wc.generate(textn_w)
print("Word Cloud for non-Duplicate Question pairs:")
plt.imshow(wc, interpolation='bilinear')
plt.axis("off")
plt.show()
```

Word Cloud for non-Duplicate Question pairs:






```
# Distribution of the token_sort_ratio
plt.figure(figsize=(10, 8))

plt.subplot(1,2,1)
sns.violinplot(x = 'is_duplicate', y = 'token_sort_ratio', data = df[0:] , )

plt.subplot(1,2,2)
sns.distplot(df[df['is_duplicate'] == 1.0]['token_sort_ratio'][0:] , label = "1", color = 'red')
sns.distplot(df[df['is_duplicate'] == 0.0]['token_sort_ratio'][0:] , label = "0" , color = 'blue' )
plt.show()
```

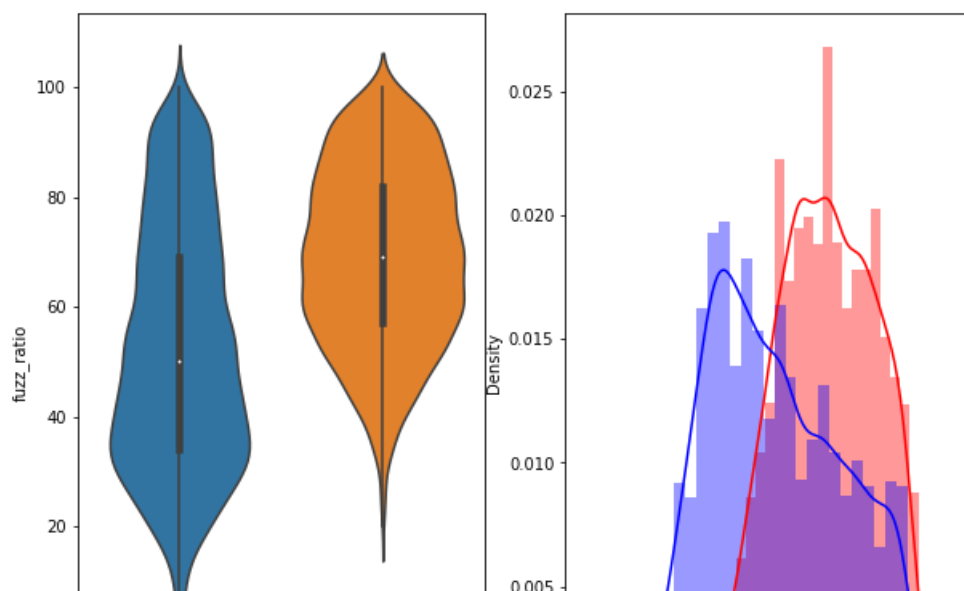



In [43]:

```
plt.figure(figsize=(10, 8))

plt.subplot(1,2,1)
sns.violinplot(x = 'is_duplicate', y = 'fuzz_ratio', data = df[0:] , )

plt.subplot(1,2,2)
sns.distplot(df[df['is_duplicate'] == 1.0]['fuzz_ratio'][0:] , label = "1", color = 'red')
sns.distplot(df[df['is_duplicate'] == 0.0]['fuzz_ratio'][0:] , label = "0" , color = 'blue' )
plt.show()
```





## 4.3 Visualization

In [44]:

```
# Using TSNE for Dimensionality reduction for 15 Features (Generated after cleaning the data) to 3
dimension

from sklearn.preprocessing import MinMaxScaler

dfp_subsampled = df[0:5000]
X = MinMaxScaler().fit_transform(dfp_subsampled[['cwc_min', 'cwc_max', 'csc_min', 'csc_max',
'ctc_min', 'ctc_max', 'last_word_eq', 'first_word_eq', 'abs_len_diff', 'mean_len', 'token_set_ratio',
'token_sort_ratio', 'fuzz_ratio', 'fuzz_partial_ratio', 'longest_substr_ratio']])
y = dfp_subsampled['is_duplicate'].values
```

In [45]:

```
tsne2d = TSNE(
    n_components=2,
    init='random', # pca
    random_state=101,
    method='barnes_hut',
    n_iter=1000,
    verbose=2,
    angle=0.5
).fit_transform(X)


[t-SNE] Computing 91 nearest neighbors...
[t-SNE] Indexed 5000 samples in 0.017s...
[t-SNE] Computed neighbors for 5000 samples in 0.382s...
[t-SNE] Computed conditional probabilities for sample 1000 / 5000
[t-SNE] Computed conditional probabilities for sample 2000 / 5000
[t-SNE] Computed conditional probabilities for sample 3000 / 5000
[t-SNE] Computed conditional probabilities for sample 4000 / 5000
[t-SNE] Computed conditional probabilities for sample 5000 / 5000
[t-SNE] Mean sigma: 0.132272
[t-SNE] Computed conditional probabilities in 0.333s
[t-SNE] Iteration 50: error = 83.1597748, gradient norm = 0.0495439 (50 iterations in 2.785s)
[t-SNE] Iteration 100: error = 70.5388184, gradient norm = 0.0098500 (50 iterations in 1.929s)
[t-SNE] Iteration 150: error = 68.5986633, gradient norm = 0.0053695 (50 iterations in 1.819s)
[t-SNE] Iteration 200: error = 67.7535858, gradient norm = 0.0038154 (50 iterations in 1.862s)
[t-SNE] Iteration 250: error = 67.2703552, gradient norm = 0.0038403 (50 iterations in 1.813s)
[t-SNE] KL divergence after 250 iterations with early exaggeration: 67.270355
[t-SNE] Iteration 300: error = 1.7787983, gradient norm = 0.0011493 (50 iterations in 1.859s)
[t-SNE] Iteration 350: error = 1.3869522, gradient norm = 0.0004775 (50 iterations in 1.863s)
[t-SNE] Iteration 400: error = 1.2232909, gradient norm = 0.0002748 (50 iterations in 1.838s)
[t-SNE] Iteration 450: error = 1.1355169, gradient norm = 0.0001849 (50 iterations in 1.833s)
[t-SNE] Iteration 500: error = 1.0818622, gradient norm = 0.0001393 (50 iterations in 1.853s)
[t-SNE] Iteration 550: error = 1.0473764, gradient norm = 0.0001151 (50 iterations in 1.875s)
[t-SNE] Iteration 600: error = 1.0237892, gradient norm = 0.0001019 (50 iterations in 1.859s)
[t-SNE] Iteration 650: error = 1.0079094, gradient norm = 0.0000901 (50 iterations in 1.832s)
[t-SNE] Iteration 700: error = 0.9969798, gradient norm = 0.0000833 (50 iterations in 1.854s)
[t-SNE] Iteration 750: error = 0.9888974, gradient norm = 0.0000865 (50 iterations in 1.858s)
[t-SNE] Iteration 800: error = 0.9828986, gradient norm = 0.0000753 (50 iterations in 1.837s)
[t-SNE] Iteration 850: error = 0.9779878, gradient norm = 0.0000710 (50 iterations in 1.890s)
[t-SNE] Iteration 900: error = 0.9736978, gradient norm = 0.0000662 (50 iterations in 1.839s)
[t-SNE] Iteration 950: error = 0.9697720, gradient norm = 0.0000648 (50 iterations in 1.867s)
[t-SNE] Iteration 1000: error = 0.9665182, gradient norm = 0.0000626 (50 iterations in 1.858s)
[t-SNE] KL divergence after 1000 iterations: 0.966518
```

In [46]:

```
df_tsne = pd.DataFrame({'x':tsne2d[:,0], 'y':tsne2d[:,1], 'label':y})

# draw the plot in appropriate place in the grid
```

```
# draw the plot in appropriate place in the grid
sns.lmplot(data= df_tsne, x='x', y='y', hue='label', fit_reg=False, size=8,palette="Set1",markers=[
's','o'])
plt.title("perplexity : {} and max_iter : {}".format(30, 1000))
plt.show()
```



## 5.Data splliting into Train and Test

In [47]:

```
import pandas as pd
import matplotlib.pyplot as plt
import re
import time
import warnings
import numpy as np
from nltk.corpus import stopwords
from sklearn.preprocessing import normalize
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.feature_extraction.text import TfidfVectorizer
warnings.filterwarnings("ignore")
import sys
import os
import pandas as pd
import numpy as np
from tqdm import tqdm

# extract word2vec vectors
# https://github.com/explosion/spaCy/issues/1721
# http://landinghub.visualstudio.com/visual-cpp-build-tools
import spacy
```

In [48]:

```
df.head(2)
```

Out[48]:

	id	qid1	qid2	question1	question2	is_duplicate	freq_qid1	freq_qid2	q1len	q2len	q1_n_words	q2_n_v
206552	206552	69616	310006	where can i find free tv shows online without ...	where can i use a credit card without the cvv ...	0	1	1	74	54	14	11
108228	108228	58375	177877	who is going to delhi this year for ias prepar...	which is the best month to go to delhi for ias...	0	1	1	52	59	10	12

In [49]:

```
# encode questions to unicode
# https://stackoverflow.com/a/6812069
# ----- python 2 -----
# df['question1'] = df['question1'].apply(lambda x: unicode(str(x), "utf-8"))
# df['question2'] = df['question2'].apply(lambda x: unicode(str(x), "utf-8"))
# ----- python 3 -----
df['question1'] = df['question1'].apply(lambda x: str(x))
df['question2'] = df['question2'].apply(lambda x: str(x))
```

In [50]:

```
df['id'].shape
```

Out[50]:

(10000,)

In [51]:

```
y = df['is_duplicate']
X = df
```

In [52]:

```
from sklearn.model_selection import train_test_split
x_tr, x_test, y_tr, y_test = train_test_split(X, y, test_size=0.3, random_state=0, stratify = y)
```

In [53]:

```
x_tr['id'].size , x_test['id'].size
```

Out[53]:

(7000, 3000)

In [54]:

```
print('-'*20+'Training Data'+ '-'*20)
print(' similar questions {}'.format(x_tr[x_tr['is_duplicate']==1]['id'].size/x_tr['id'].size))
print(' not similar questions {}'.format(x_tr[x_tr['is_duplicate']==0]['id'].size/x_tr['id'].size))

print('-'*20+'Test Data'+ '-'*20)
print(' similar questions are {}'.format(x_test[x_test['is_duplicate']==1]['id'].size/x_test['id'].size))
print(' not similar questions are {}'.format(x_test[x_test['is_duplicate']==0]['id'].size/x_test['id'].size))
```

```

-----Training Data-----
similar questions 0.37185714285714283
not similar questions 0.6281428571428571
-----Test Data-----
similar questions are 0.372
not similar questions are 0.628

```

In [55]:

```
x_tr.head(2)
```

Out[55]:

	id	qid1	qid2	question1	question2	is_duplicate	freq_qid1	freq_qid2	q1len	q2len	q1_n_words	q2_n
95094	95094	158674	158675	is the liberal democratic party truly liberal	what does the liberal democratic party believe...	0	1	1	46	50	7	8
220488	220488	327589	327590	what are some fun bet ideas	what are fun stakes to set when making a bet w...	0	1	1	28	63	6	13

## 6. Featurizing text data with tfidf weighted word-vectors

### 6.1 on question\_1

In [56]:

```

from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.feature_extraction.text import CountVectorizer
# merge texts
que_tr_1 = list(x_tr['question1'])
que_test_1 = list(x_test['question1'])

tfidf = TfidfVectorizer(lowercase=False, max_features = 100)
tfidf.fit(que_tr_1)

```

Out[56]:

```

TfidfVectorizer(analyzer='word', binary=False, decode_error='strict',
dtype=<class 'numpy.float64'>, encoding='utf-8',
input='content', lowercase=False, max_df=1.0, max_features=100,
min_df=1, ngram_range=(1, 1), norm='l2', preprocessor=None,
smooth_idf=True, stop_words=None, strip_accents=None,
sublinear_tf=False, token_pattern='(?u)\\b\\w\\w+\\b',
tokenizer=None, use_idf=True, vocabulary=None)

```

In [57]:

```

x_tr_que_tfidf_1 = tfidf.transform(que_tr_1)
x_test_que_tfidf_1 = tfidf.transform(que_test_1)

```

In [58]:

```

print("Shape of matrix after one hot encodig ",x_tr_que_tfidf_1.shape)
print("Shape of matrix after one hot encodig ",x_test_que_tfidf_1.shape)

```

Shape of matrix after one hot encodig (7000, 100)

Shape of matrix after one hot encoding (3000, 100)

- After we find TF-IDF scores, we convert each question to a weighted average of word2vec vectors by these scores.
- here we use a pre-trained GLOVE model which comes free with "Spacy". <https://spacy.io/usage/vectors-similarity>
- It is trained on Wikipedia and therefore, it is stronger in terms of word semantics.

In [59]:

```
df['id'].shape
```

Out[59]:

```
(10000,)
```

In [60]:

```
import spacy
nlp = spacy.load('en_core_web_sm') # en_vectors_web_lg, which includes over 1 million unique
vectors.
```

In [61]:

```
# dict key:word and value:tf-idf score
word2tfidf = dict(zip(tfidf.get_feature_names(), tfidf.idf_))
```

In [62]:

```
vecs1_tr = []
# https://github.com/noamraph/tqdm
# tqdm is used to print the progress bar
for qul in tqdm(list(x_tr['question1'])):
    doc1 = nlp(qul)
    # 384 is the number of dimensions of vectors
    n = len(doc1)
    if(n!=0):
        m = len(doc1[0].vector)
    else:
        m = 96
    mean_vec1 = np.zeros([n, m])
    for word1 in doc1:
        # word2vec
        vec1 = word1.vector
        # fetch df score
        try:
            idf = word2tfidf[str(word1)]
        except:
            idf = 0
        # compute final vec
        mean_vec1 += vec1 * idf
    mean_vec1 = mean_vec1.mean(axis=0)
    vecs1_tr.append(mean_vec1)
```

100%|██████████| 7000/7000 [03:56<00:00, 29.58it/s]

In [63]:

```
vecs1_tr[0].shape, vecs1_tr[100].shape, len(doc1), len(doc1[0].vector)
```

Out[63]:

```
((96,), (96,), 7, 96)
```

In [64]:

```
doc1[0].vector
```

Out[64]:

```
array([-1.7159698, -0.94028306, 1.4302827, -4.8813066, 5.602137, ...,
```

```

0.41269925, -2.3720655 , 1.1801261 , -2.6673112 , 0.47985882,
0.0292834 , -0.5730588 , -1.4184704 , 1.4368011 , 4.124066 ,
-0.11633208, 1.841193 , 2.8167768 , 1.6611824 , 1.0557799 ,
-3.7842743 , 0.60988855, -1.2097673 , -1.5597488 , -4.5673475 ,
-0.53190565, -4.439723 , -2.4376507 , -0.4086672 , 1.2029594 ,
0.15232885, 4.910077 , -1.9216218 , 2.8664198 , 0.74302423,
-2.1705668 , 0.13028345, 0.11531692, 3.9189754 , 0.15966076,
-1.416482 , 1.2364417 , -0.93643475, -2.3468156 , 4.6074243 ,
-1.3140641 , -0.5752812 , -1.7728248 , -1.2327435 , 0.7378238 ,
1.7125928 , 0.04458308, -1.7939475 , -1.401803 , 2.467896 ,
0.18753523, 0.07729542, -2.5627563 , -0.61417913, -0.76402867,
1.3572025 , -0.1970613 , -1.8572826 , -0.22799663, -2.5539856 ,
6.2849994 , 0.28354916, -2.2318506 , 1.4054313 , -0.16206911,
0.72770584, -0.54049134, 1.4264905 , 0.9644554 , -4.271908 ,
1.7764267 , 1.0350004 , -0.6052711 , 2.7613888 , -0.70816153,
0.18992624, -0.656117 , -0.67705107, 2.6094098 , 4.0956664 ,
-0.882318 , -2.0373385 , 1.888547 , 0.5793896 , 3.584142 ,
-3.0542512 , -2.5895586 , -0.3792294 , 0.20076397, -0.12577087,
-0.10814381], dtype=float32)

```

In [65]:

```
type(vecs1_tr)
```

Out[65]:

list

In [66]:

```
x_tr['tfidf_w2v_q1'] = list(vecs1_tr)
```

In [67]:

```
x_tr.head(1)
```

Out[67]:

	id	qid1	qid2	question1	question2	is_duplicate	freq_qid1	freq_qid2	q1len	q2len	q1_n_words	q2_n_w
95094	95094	158674	158675	is the liberal democratic party truly liberal	what does the liberal democratic party believe...	0	1	1	46	50	7	8

In [68]:

```

vecs1_test = []
# https://github.com/noamraph/tqdm
# tqdm is used to print the progress bar
for qul in tqdm(list(x_test['question1'])):
    doc1 = nlp(qul)
    # 384 is the number of dimensions of vectors
    n = len(doc1)
    if (n!=0):
        m = len(doc1[0].vector)
    else:
        m = 96
    mean_vec1 = np.zeros([n, m])
    for word1 in doc1:
        # word2vec
        vec1 = word1.vector
        # fetch df score
        try:
            idf = word2tfidf[str(word1)]
        except:
            idf = 0
        # compute final vec
        mean_vec1 += vec1 * idf

```



```
mean_vec1 = mean_vec1.mean(axis=0)
vecs1_test.append(mean_vec1)
```

100%|██████████| 3000/3000 [01:43<00:00, 29.03it/s]

In [69]:

```
x_test['tfidf_w2v_q1'] = list(vecs1_test)
```

In [70]:

```
vecs1_test[0].shape
```

Out[70]:

(96,)

In [71]:

```
x_test.head(1)
```

Out[71]:

	id	qid1	qid2	question1	question2	is_duplicate	freq_qid1	freq_qid2	q1len	q2len	q1_n_words	q2_n_v
144592	144592	74372	228794	why does human life exist on earth	is there a place on earth where no life exists	0	1	1	35	47	7	10

## 6.2 on question\_2

In [72]:

```
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.feature_extraction.text import CountVectorizer
# merge texts
que_tr_2 = list(x_tr['question2'])
que_test_2 = list(x_test['question2'])

tfidf = TfidfVectorizer(lowercase=False, max_features = 100)
tfidf.fit(que_tr_2)
```

Out[72]:

```
TfidfVectorizer(analyzer='word', binary=False, decode_error='strict',
dtype=<class 'numpy.float64'>, encoding='utf-8',
input='content', lowercase=False, max_df=1.0, max_features=100,
min_df=1, ngram_range=(1, 1), norm='l2', preprocessor=None,
smooth_idf=True, stop_words=None, strip_accents=None,
sublinear_tf=False, token_pattern='(?u)\\b\\w\\w+\\b',
tokenizer=None, use_idf=True, vocabulary=None)
```

In [73]:

```
x_tr_que_tfidf_2 = tfidf.transform(que_tr_2)

x_test_que_tfidf_2 = tfidf.transform(que_test_2)
```

In [74]:

```
print("Shape of matrix after one hot encodig ",x_tr_que_tfidf_1.shape)
print("Shape of matrix after one hot encodig ",x_test_que_tfidf_1.shape)
```

```
Shape of matrix after one hot encodig (7000, 100)
Shape of matrix after one hot encodig (3000, 100)
```

- After we find TF-IDF scores, we convert each question to a weighted average of word2vec vectors by these scores.
- here we use a pre-trained GLOVE model which comes free with "Spacy". <https://spacy.io/usage/vectors-similarity>
- It is trained on Wikipedia and therefore, it is stronger in terms of word semantics.

```
In [75]:
```

```
df['id'].shape
```

```
Out[75]:
```

```
(10000,)
```

```
In [76]:
```

```
import spacy
nlp = spacy.load('en_core_web_sm') # en_vectors_web_lg, which includes over 1 million unique
vectors.
```

```
In [77]:
```

```
# dict key:word and value:tf-idf score
word2tfidf = dict(zip(tfidf.get_feature_names(), tfidf.idf_))
```

```
In [78]:
```

```
vecs2_tr = []
for qu2 in tqdm(list(x_tr['question2'])):
    doc2 = nlp(qu2)
    n = len(doc2)
    if n!=0:
        m = len(doc2[0].vector)
    else:
        m = 96
    mean_vec2 = np.zeros([n, m])
    for word2 in doc2:
        # word2vec
        vec2 = word2.vector
        # fetch df score
        try:
            idf = word2tfidf[str(word2)]
        except:
            #print word
            idf = 0
        # compute final vec
        mean_vec2 += vec2 * idf
    mean_vec2 = mean_vec2.mean(axis=0)
    vecs2_tr.append(mean_vec2)
```

```
100%|██████████| 7000/7000 [04:04<00:00, 28.66it/s]
```

```
In [79]:
```

```
x_tr['tfidf_w2v_q2'] = list(vecs2_tr)
```

```
In [80]:
```

```
x_tr.head(1)
```

```
Out[80]:
```

	id	qid1	qid2	question1	question2	is_duplicate	freq_qid1	freq_qid2	q1len	q2len	q1_n_words	q2_n_w
				is the	what does							

95094	95094	158674	158672	liberal question1 party truly liberal	the liberal question2 party believe...	is_duplicate	freq_qid1	freq_qid2	q1len	q2len	q1_n_words	q2_n_w

In [81]:

```
vecs2_test = []
for qu2 in tqdm(list(x_test['question2'])):
    doc2 = nlp(qu2)
    n = len(doc2)
    if (n!=0):
        m = len(doc2[0].vector)
    else:
        m = 96
    mean_vec2 = np.zeros([n, m])
    for word2 in doc2:
        # word2vec
        vec2 = word2.vector
        # fetch df score
        try:
            idf = word2tfidf[str(word2)]
        except:
            #print word
            idf = 0
        # compute final vec
        mean_vec2 += vec2 * idf
    mean_vec2 = mean_vec2.mean(axis=0)
    vecs2_test.append(mean_vec2)
```

100%|██████████| 3000/3000 [01:43<00:00, 29.08it/s]

In [82]:

```
x_test['tfidf_w2v_q2'] = list(vecs2_test)
```

In [83]:

```
x_test.head(1)
```

Out[83]:

	id	qid1	qid2	question1	question2	is_duplicate	freq_qid1	freq_qid2	q1len	q2len	q1_n_words	q2_n_v
144592	144592	74372	228794	why does human life exist on earth	is there a place on earth where no life exists	0	1	1	35	47	7	10

In [84]:

```
x_tr.to_csv(file_path+"x_tr.csv")
x_test.to_csv(file_path+"x_test.csv")
```

## 7.Creating Final Features dataframe

### 7.1 TFIDF-W2V

In [85]:

```
# #prepro_features_train.csv (Simple Preprocessing Feartures)
# #nlp_features_train.csv (NLP Features)
# if os.path.isfile('nlp_features_train.csv'):
```

```
# dfnlp = pd.read_csv("nlp_features_train.csv",encoding='latin-1')
# else:
#     print("download nlp_features_train.csv from drive or run previous notebook")

# if os.path.isfile('df_fe_without_preprocessing_train.csv'):
#     dfppro = pd.read_csv("df_fe_without_preprocessing_train.csv",encoding='latin-1')
# else:
#     print("download df_fe_without_preprocessing_train.csv from drive or run previous notebook")
```

In [86]:

```
# df1 = dfnlp.drop(['qid1','qid2','question1','question2'],axis=1)
# df2 = dfppro.drop(['qid1','qid2','question1','question2','is_duplicate'],axis=1)
# df3 = df.drop(['qid1','qid2','question1','question2','is_duplicate'],axis=1)
x_tr_tfidf_w2v_q1 = pd.DataFrame(x_tr.tfidf_w2v_q1.values.tolist(), index= x_tr.index)
x_tr_tfidf_w2v_q2 = pd.DataFrame(x_tr.tfidf_w2v_q2.values.tolist(), index= x_tr.index)
```

In [87]:

```
x_tr_tfidf_w2v_q1.shape, x_tr_tfidf_w2v_q2.shape
```

Out[87]:

```
((7000, 96), (7000, 96))
```

In [88]:

```
x_test_tfidf_w2v_q1 = pd.DataFrame(x_test.tfidf_w2v_q1.values.tolist(), index= x_test.index)
x_test_tfidf_w2v_q2 = pd.DataFrame(x_test.tfidf_w2v_q2.values.tolist(), index= x_test.index)
```

In [89]:

```
x_test_tfidf_w2v_q1.shape, x_test_tfidf_w2v_q2.shape
```

Out[89]:

```
((3000, 96), (3000, 96))
```

In [90]:

```
x_tr_tfidf_w2v_q1['id'] = x_tr['id']
x_tr_tfidf_w2v_q2['id'] = x_tr['id']
```

In [91]:

```
x_test_tfidf_w2v_q1['id'] = x_test['id']
x_test_tfidf_w2v_q2['id'] = x_test['id']
```

In [92]:

```
x_test_tfidf_w2v_q2.head(2)
```

Out[92]:

	0	1	2	3	4	5	6	7	8	9	
<b>144592</b>	32.704681	4.057020	52.485193	-29.174418	29.249085	-8.565338	21.239076	20.250900	-27.405324	44.485319	3.97
<b>119904</b>	-4.982040	-11.763009	6.502761	-17.936214	23.026928	-13.473953	19.332658	-3.386513	-15.488996	13.496412	-2.88

2 rows × 97 columns



In [93]:

```
x_tr.head(2)
```

Out [93]:

	id	qid1	qid2	question1	question2	is_duplicate	freq_qid1	freq_qid2	q1len	q2len	q1_n_words	q2_n
95094	95094	158674	158675	is the liberal democratic party truly liberal	what does the liberal democratic party believe...	0	1	1	46	50	7	8
220488	220488	327589	327590	what are some fun bet ideas	what are fun stakes to set when making a bet w...	0	1	1	28	63	6	13

In [94]:

```
x_test.head(2)
```

Out [94]:

	id	qid1	qid2	question1	question2	is_duplicate	freq_qid1	freq_qid2	q1len	q2len	q1_n_words	q2_n
144592	144592	74372	228794	why does human life exist on earth	is there a place on earth where no life exists	0	1	1	35	47	7	10
119904	119904	194577	194578	what are the worst things about bollywood	what is the worst thing about bollywood movies	1	1	1	42	47	7	8

In [95]:

```
x_tr.shape, x_test.shape
```

Out [95]:

```
((7000, 34), (3000, 34))
```

In [96]:

```
if not os.path.isfile(file_path + 'final_features_tfidf_w2v_tr.csv'):  
  
    print(x_tr.shape)  
    x_tr = x_tr.merge(x_tr_tfidf_w2v_q1, on='id',how='left')  
    x_tr = x_tr.merge(x_tr_tfidf_w2v_q2, on='id',how='left')  
    print(x_tr.shape)  
  
    print(x_test.shape)  
    x_test = x_test.merge(x_test_tfidf_w2v_q1, on='id',how='left')  
    x_test = x_test.merge(x_test_tfidf_w2v_q2, on='id',how='left')  
    print(x_test.shape)  
  
    x_tr.to_csv(file_path + 'final_features_tfidf_w2v_tr.csv')  
    x_test.to_csv(file_path + 'final_features_tfidf_w2v_test.csv')
```

```
(7000, 34)
```

```
(7000, 226)
(3000, 34)
(3000, 226)
```

In [97]:

```
x_tr.head(2)
```

Out[97]:

	id	qid1	qid2	question1	question2	is_duplicate	freq_qid1	freq_qid2	q1len	q2len	q1_n_words	q2_n_word
0	95094	158674	158675	is the liberal democratic party truly liberal	what does the liberal democratic party believe...	0	1	1	46	50	7	8
1	220488	327589	327590	what are some fun bet ideas	what are fun stakes to set when making a bet w...	0	1	1	28	63	6	13

2 rows × 226 columns



In [98]:

```
x_test.head(2)
```

Out[98]:

	id	qid1	qid2	question1	question2	is_duplicate	freq_qid1	freq_qid2	q1len	q2len	q1_n_words	q2_n_word:
0	144592	74372	228794	why does human life exist on earth	is there a place on earth where no life exists	0	1	1	35	47	7	10
1	119904	194577	194578	what are the worst things about bollywood	what is the worst thing about bollywood movies	1	1	1	42	47	7	8

2 rows × 226 columns



## 7.2 TFIDF

In [99]:

```
x_tr.columns.values
```

Out[99]:

```
array(['id', 'qid1', 'qid2', 'question1', 'question2', 'is_duplicate',
       'freq_qid1', 'freq_qid2', 'q1len', 'q2len', 'q1_n_words',
       'q2_n_words', 'word_Common', 'word_Total', 'word_share',
       'freq_q1-q2', 'freq_q1-q2', 'cwc_min', 'cwc_max', 'csc_min',
       'csc_max', 'ctc_min', 'ctc_max', 'last_word_eq', 'first_word_eq',
       'abs_len_diff', 'mean_len', 'token_set_ratio', 'token_sort_ratio',
```

```
'fuzz_ratio', 'fuzz_partial_ratio', 'longest_substr_ratio',
'tfidf_w2v_q1', 'tfidf_w2v_q2', '0_x', '1_x', '2_x', '3_x', '4_x',
'5_x', '6_x', '7_x', '8_x', '9_x', '10_x', '11_x', '12_x', '13_x',
'14_x', '15_x', '16_x', '17_x', '18_x', '19_x', '20_x', '21_x',
'22_x', '23_x', '24_x', '25_x', '26_x', '27_x', '28_x', '29_x',
'30_x', '31_x', '32_x', '33_x', '34_x', '35_x', '36_x', '37_x',
'38_x', '39_x', '40_x', '41_x', '42_x', '43_x', '44_x', '45_x',
'46_x', '47_x', '48_x', '49_x', '50_x', '51_x', '52_x', '53_x',
'54_x', '55_x', '56_x', '57_x', '58_x', '59_x', '60_x', '61_x',
'62_x', '63_x', '64_x', '65_x', '66_x', '67_x', '68_x', '69_x',
'70_x', '71_x', '72_x', '73_x', '74_x', '75_x', '76_x', '77_x',
'78_x', '79_x', '80_x', '81_x', '82_x', '83_x', '84_x', '85_x',
'86_x', '87_x', '88_x', '89_x', '90_x', '91_x', '92_x', '93_x',
'94_x', '95_x', '0_y', '1_y', '2_y', '3_y', '4_y', '5_y', '6_y',
'7_y', '8_y', '9_y', '10_y', '11_y', '12_y', '13_y', '14_y',
'15_y', '16_y', '17_y', '18_y', '19_y', '20_y', '21_y', '22_y',
'23_y', '24_y', '25_y', '26_y', '27_y', '28_y', '29_y', '30_y',
'31_y', '32_y', '33_y', '34_y', '35_y', '36_y', '37_y', '38_y',
'39_y', '40_y', '41_y', '42_y', '43_y', '44_y', '45_y', '46_y',
'47_y', '48_y', '49_y', '50_y', '51_y', '52_y', '53_y', '54_y',
'55_y', '56_y', '57_y', '58_y', '59_y', '60_y', '61_y', '62_y',
'63_y', '64_y', '65_y', '66_y', '67_y', '68_y', '69_y', '70_y',
'71_y', '72_y', '73_y', '74_y', '75_y', '76_y', '77_y', '78_y',
'79_y', '80_y', '81_y', '82_y', '83_y', '84_y', '85_y', '86_y',
'87_y', '88_y', '89_y', '90_y', '91_y', '92_y', '93_y', '94_y',
'95_y'], dtype=object)
```

In [100]:

```
x_tr = x_tr[['id', 'qid1', 'qid2', 'question1', 'question2', 'is_duplicate',
'freq_qid1', 'freq_qid2', 'qlen', 'q2len', 'q1_n_words',
'q2_n_words', 'word_Common', 'word_Total', 'word_share',
'freq_q1+q2', 'freq_q1-q2', 'cwc_min', 'cwc_max', 'csc_min',
'csc_max', 'ctc_min', 'ctc_max', 'last_word_eq', 'first_word_eq',
'abs_len_diff', 'mean_len', 'token_set_ratio', 'token_sort_ratio',
'fuzz_ratio', 'fuzz_partial_ratio', 'longest_substr_ratio']]
```

In [101]:

```
x_test = x_test[['id', 'qid1', 'qid2', 'question1', 'question2', 'is_duplicate',
'freq_qid1', 'freq_qid2', 'qlen', 'q2len', 'q1_n_words',
'q2_n_words', 'word_Common', 'word_Total', 'word_share',
'freq_q1+q2', 'freq_q1-q2', 'cwc_min', 'cwc_max', 'csc_min',
'csc_max', 'ctc_min', 'ctc_max', 'last_word_eq', 'first_word_eq',
'abs_len_diff', 'mean_len', 'token_set_ratio', 'token_sort_ratio',
'fuzz_ratio', 'fuzz_partial_ratio', 'longest_substr_ratio']]
```

In [102]:

```
x_tr.head(2)
```

Out[102]:

	id	qid1	qid2	question1	question2	is_duplicate	freq_qid1	freq_qid2	q1len	q2len	q1_n_words	q2_n_word
0	95094	158674	158675	is the liberal democratic party truly liberal	what does the liberal democratic party believe...	0	1	1	46	50	7	8
1	220488	327589	327590	what are some fun bet ideas	what are fun stakes to set when making a bet w...	0	1	1	28	63	6	13

In [103]:

```
x_tr_que_tfidf_1.shape, x_tr_que_tfidf_2.shape, x_test_que_tfidf_1.shape, x_test_que_tfidf_2.shape
```

Out[103]:

```
((7000, 100), (7000, 100), (3000, 100), (3000, 100))
```

In [104]:

```
# df1 = dfnlp.drop(['qid1','qid2','question1','question2'],axis=1)
# df2 = dfppro.drop(['qid1','qid2','question1','question2','is_duplicate'],axis=1)
# df3 = df.drop(['qid1','qid2','question1','question2','is_duplicate'],axis=1)
x_tr_que_tfidf_1 = pd.DataFrame(x_tr_que_tfidf_1.todense())
x_tr_que_tfidf_2 = pd.DataFrame(x_tr_que_tfidf_2.todense())

x_test_que_tfidf_1 = pd.DataFrame(x_test_que_tfidf_1.todense())
x_test_que_tfidf_2 = pd.DataFrame(x_test_que_tfidf_2.todense())
```

In [105]:

```
x_tr_que_tfidf_1.shape, x_tr_que_tfidf_2.shape, x_test_que_tfidf_1.shape, x_test_que_tfidf_2.shape
```

Out[105]:

```
((7000, 100), (7000, 100), (3000, 100), (3000, 100))
```

In [106]:

```
x_tr_que_tfidf_1['id'] = x_tr['id']
x_tr_que_tfidf_2['id'] = x_tr['id']
```

In [107]:

```
x_test_que_tfidf_1['id'] = x_test['id']
x_test_que_tfidf_2['id'] = x_test['id']
```

In [108]:

```
x_tr_que_tfidf_1.head(2)
```

Out[108]:

	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27
0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.000000	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
1	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.517875	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0

2 rows × 101 columns



In [109]:

```
x_test_que_tfidf_2.head(2)
```

Out[109]:

	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27
0	0.000000	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
1	0.809968	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0

2 rows × 101 columns



In [110]:

```
x_tr.shape, x_test.shape
```



$((7000, 32), (3000, 32))$ 

```
if not os.path.isfile(file_path + 'final_features_tfidf_tr.csv'):

    print(x_tr.shape)
    x_tr = x_tr.merge(x_tr_que_tfidf_1, on='id', how='left')
    x_tr = x_tr.merge(x_tr_que_tfidf_2, on='id', how='left')
    print(x_tr.shape)

    print(x_test.shape)
    x_test = x_test.merge(x_test_que_tfidf_1, on='id', how='left')
    x_test = x_test.merge(x_test_que_tfidf_2, on='id', how='left')
    print(x_test.shape)

x_tr.to_csv(file_path + 'final_features_tfidf_tr.csv')
x_test.to_csv(file_path + 'final_features_tfidf_test.csv')
```

In [112]:

Out [112] :

2 rows × 232 columns

```
x.tr.isna().any().values
```

[illegible]

In [114]:

Out[114]:

2 rows × 232 columns



```
import pandas as pd
import matplotlib.pyplot as plt
import re
import time
import warnings
import sqlite3
from sqlalchemy import create_engine # database connection
import csv
import os
warnings.filterwarnings("ignore")
import datetime as dt
import numpy as np
from nltk.corpus import stopwords
from sklearn.decomposition import TruncatedSVD
from sklearn.preprocessing import normalize
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.manifold import TSNE
import seaborn as sns
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import confusion_matrix
from sklearn.metrics.classification import accuracy_score, log_loss
from sklearn.feature_extraction.text import TfidfVectorizer
from collections import Counter
from scipy.sparse import hstack
from sklearn.multiclass import OneVsRestClassifier
from sklearn.svm import SVC
from sklearn.model_selection import StratifiedKFold
```

```

from collections import Counter, defaultdict
from sklearn.calibration import CalibratedClassifierCV
from sklearn.naive_bayes import MultinomialNB
from sklearn.naive_bayes import GaussianNB
from sklearn.model_selection import train_test_split
from sklearn.model_selection import GridSearchCV
import math
from sklearn.metrics import normalized_mutual_info_score
from sklearn.ensemble import RandomForestClassifier

from sklearn.model_selection import cross_val_score
from sklearn.linear_model import SGDClassifier
from mlxtend.classifier import StackingClassifier

from sklearn import model_selection
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import precision_recall_curve, auc, roc_curve

import time
import xgboost as xgb
from sklearn.model_selection import RandomizedSearchCV

```

## 8.1 TFIDF - W2V

### 8.1.1 Train data

In [116]:

```

#http://www.sqlitetutorial.net/sqlite-python/create-tables/
def create_connection(db_file):
    """ create a database connection to the SQLite database
        specified by db_file
    :param db_file: database file
    :return: Connection object or None
    """
    try:
        conn = sqlite3.connect(db_file)
        return conn
    except Error as e:
        print(e)

    return None

def checkTableExists(dbcon):
    cursr = dbcon.cursor()
    str = "select name from sqlite_master where type='table'"
    table_names = cursr.execute(str)
    print("Tables in the database:")
    tables = table_names.fetchall()
    print(tables[0][0])
    print(tables)
    return(len(tables))

```

In [117]:

```

#Creating db file from csv
import sqlalchemy
if not os.path.isfile(file_path + 'tfidf_w2v_tr.db'):
    print("not present in drive")
    disk_engine = sqlalchemy.create_engine('sqlite:///'+file_path+'tfidf_w2v_tr.db')
    start = dt.datetime.now()
    chunksize = 180000
    j = 0
    index_start = 1
    for df in pd.read_csv(file_path+'final_features_tfidf_w2v_tr.csv',
                        names=['id', 'qid1', 'qid2', 'question1', 'question2', 'is_duplicate',
                              'freq_qid1', 'freq_qid2', 'q1len', 'q2len', 'q1_n_words',
                              'q2_n_words', 'word_Common', 'word_Total', 'word_share',
                              'freq q1+q2', 'freq q1-q2', 'cwc min', 'cwc max', 'csc min',

```

```

'csc_max', 'ctc_min', 'ctc_max', 'last_word_eq', 'first_word_eq',
'abs_len_diff', 'mean_len', 'token_set_ratio', 'token_sort_ratio',
'fuzz_ratio', 'fuzz_partial_ratio', 'longest_substr_ratio',
'tfidf_w2v_q1', 'tfidf_w2v_q2', '0_x', '1_x', '2_x', '3_x', '4_x',
'5_x', '6_x', '7_x', '8_x', '9_x', '10_x', '11_x', '12_x', '13_x',
'14_x', '15_x', '16_x', '17_x', '18_x', '19_x', '20_x', '21_x',
'22_x', '23_x', '24_x', '25_x', '26_x', '27_x', '28_x', '29_x',
'30_x', '31_x', '32_x', '33_x', '34_x', '35_x', '36_x', '37_x',
'38_x', '39_x', '40_x', '41_x', '42_x', '43_x', '44_x', '45_x',
'46_x', '47_x', '48_x', '49_x', '50_x', '51_x', '52_x', '53_x',
'54_x', '55_x', '56_x', '57_x', '58_x', '59_x', '60_x', '61_x',
'62_x', '63_x', '64_x', '65_x', '66_x', '67_x', '68_x', '69_x',
'70_x', '71_x', '72_x', '73_x', '74_x', '75_x', '76_x', '77_x',
'78_x', '79_x', '80_x', '81_x', '82_x', '83_x', '84_x', '85_x',
'86_x', '87_x', '88_x', '89_x', '90_x', '91_x', '92_x', '93_x',
'94_x', '95_x', '0_y', '1_y', '2_y', '3_y', '4_y', '5_y', '6_y',
'7_y', '8_y', '9_y', '10_y', '11_y', '12_y', '13_y', '14_y',
'15_y', '16_y', '17_y', '18_y', '19_y', '20_y', '21_y', '22_y',
'23_y', '24_y', '25_y', '26_y', '27_y', '28_y', '29_y', '30_y',
'31_y', '32_y', '33_y', '34_y', '35_y', '36_y', '37_y', '38_y',
'39_y', '40_y', '41_y', '42_y', '43_y', '44_y', '45_y', '46_y',
'47_y', '48_y', '49_y', '50_y', '51_y', '52_y', '53_y', '54_y',
'55_y', '56_y', '57_y', '58_y', '59_y', '60_y', '61_y', '62_y',
'63_y', '64_y', '65_y', '66_y', '67_y', '68_y', '69_y', '70_y',
'71_y', '72_y', '73_y', '74_y', '75_y', '76_y', '77_y', '78_y',
'79_y', '80_y', '81_y', '82_y', '83_y', '84_y', '85_y', '86_y',
'87_y', '88_y', '89_y', '90_y', '91_y', '92_y', '93_y', '94_y',
'95_y'],
        chunksize=chunksize, iterator=True, encoding='utf-8', ):
df.index += index_start
j+=1
print('{} rows'.format(j*chunksize))
df.to_sql('data', disk_engine, if_exists='append')
index_start = df.index[-1] + 1

```

not present in drive  
180000 rows

In [118]:

```

read_db = file_path + 'tfidf_w2v_tr.db'
conn_r = create_connection(read_db)
checkTableExists(conn_r)
conn_r.close()

```

Tables in the database:  
data  
[('data',)]

In [119]:

```

# try to sample data according to the computing power you have
if os.path.isfile(read_db):
    conn_r = create_connection(read_db)
    if conn_r is not None:
        # for selecting first 1M rows
        # data = pd.read_sql_query("""SELECT * FROM data LIMIT 100001;""", conn_r)
        data = pd.read_sql_query("""SELECT * From data ORDER BY RANDOM() LIMIT 100001;""", conn_r)

        # for selecting random points
        # data = pd.read_sql_query("SELECT * From data ORDER BY RANDOM() LIMIT 100001;", conn_r)
    conn_r.commit()
    conn_r.close()

```

In [120]:

```
data.shape
```

Out[120]:

(7001, 227)

In [121]:

```
data[data['is_duplicate']=='is_duplicate']
```

Out[121]:

	index	id	qid1	qid2	question1	question2	is_duplicate	freq_qid1	freq_qid2	q1len	q2len	q1_n_words	q2_n_words
1094	NaN	id	qid1	qid2	question1	question2	is_duplicate	freq_qid1	freq_qid2	q1len	q2len	q1_n_words	q2_n_words

1 rows × 227 columns

In [122]:

```
data = data.drop(data.index[data['is_duplicate']=='is_duplicate'])
```

In [123]:

```
data.head(2)
```

Out[123]:

	index	id	qid1	qid2	question1	question2	is_duplicate	freq_qid1	freq_qid2	q1len	q2len	q1_n_words	q2_n_
0	6319.0	98977	86348	31174	what are some good lyric prank songs for your ...	what is good song for a best friend lyric prank	1	1	1	58	47	11	9
1	4854.0	216800	3066	26378	how can i improve my writing skills for writin...	how do i improve my english writing and speaki...	1	1	1	55	56	11	10

2 rows × 227 columns

In [124]:

```
data[pd.isnull(data).any(axis=1)]
```

Out[124]:

	index	id	qid1	qid2	question1	question2	is_duplicate	freq_qid1	freq_qid2	q1len	q2len	q1_n_words	q2_n_words	wo
--	-------	----	------	------	-----------	-----------	--------------	-----------	-----------	-------	-------	------------	------------	----

0 rows × 227 columns

In [125]:

```
# remove the first row
# data.drop(data.index[0], inplace=True)
y_train = data['is_duplicate'].values
data.drop(['id', 'index', 'is_duplicate'], axis=1, inplace=True)
```

In [126]:

```
data.shape, y_train.shape
```

Out[126]:

((7000, 224), (7000,))

In [127]:

```
data.head(2)
```

Out[127]:

	qid1	qid2	question1	question2	freq_qid1	freq_qid2	q1len	q2len	q1_n_words	q2_n_words	word_Common	word_
0	86348	31174	what are some good lyric prank songs for your ...	what is good song for a best friend lyric prank	1	1	58	47	11	9	4.0	20.0
1	3066	26378	how can i improve my writing skills for writin...	how do i improve my english writing and speaki...	1	1	55	56	11	10	5.0	20.0

2 rows × 224 columns



In [128]:

```
data.columns.values
```

Out[128]:

```
array(['qid1', 'qid2', 'question1', 'question2', 'freq_qid1', 'freq_qid2',
      'q1len', 'q2len', 'q1_n_words', 'q2_n_words', 'word_Common',
      'word_Total', 'word_share', 'freq_q1+q2', 'freq_q1-q2', 'cwc_min',
      'cwc_max', 'csc_min', 'csc_max', 'ctc_min', 'ctc_max',
      'last_word_eq', 'first_word_eq', 'abs_len_diff', 'mean_len',
      'token_set_ratio', 'token_sort_ratio', 'fuzz_ratio',
      'fuzz_partial_ratio', 'longest_substr_ratio', 'tfidf_w2v_q1',
      'tfidf_w2v_q2', '0_x', '1_x', '2_x', '3_x', '4_x', '5_x', '6_x',
      '7_x', '8_x', '9_x', '10_x', '11_x', '12_x', '13_x', '14_x',
      '15_x', '16_x', '17_x', '18_x', '19_x', '20_x', '21_x', '22_x',
      '23_x', '24_x', '25_x', '26_x', '27_x', '28_x', '29_x', '30_x',
      '31_x', '32_x', '33_x', '34_x', '35_x', '36_x', '37_x', '38_x',
      '39_x', '40_x', '41_x', '42_x', '43_x', '44_x', '45_x', '46_x',
      '47_x', '48_x', '49_x', '50_x', '51_x', '52_x', '53_x', '54_x',
      '55_x', '56_x', '57_x', '58_x', '59_x', '60_x', '61_x', '62_x',
      '63_x', '64_x', '65_x', '66_x', '67_x', '68_x', '69_x', '70_x',
      '71_x', '72_x', '73_x', '74_x', '75_x', '76_x', '77_x', '78_x',
      '79_x', '80_x', '81_x', '82_x', '83_x', '84_x', '85_x', '86_x',
      '87_x', '88_x', '89_x', '90_x', '91_x', '92_x', '93_x', '94_x',
      '95_x', '0_y', '1_y', '2_y', '3_y', '4_y', '5_y', '6_y', '7_y',
      '8_y', '9_y', '10_y', '11_y', '12_y', '13_y', '14_y', '15_y',
      '16_y', '17_y', '18_y', '19_y', '20_y', '21_y', '22_y', '23_y',
      '24_y', '25_y', '26_y', '27_y', '28_y', '29_y', '30_y', '31_y',
      '32_y', '33_y', '34_y', '35_y', '36_y', '37_y', '38_y', '39_y',
      '40_y', '41_y', '42_y', '43_y', '44_y', '45_y', '46_y', '47_y',
      '48_y', '49_y', '50_y', '51_y', '52_y', '53_y', '54_y', '55_y',
      '56_y', '57_y', '58_y', '59_y', '60_y', '61_y', '62_y', '63_y',
      '64_y', '65_y', '66_y', '67_y', '68_y', '69_y', '70_y', '71_y',
      '72_y', '73_y', '74_y', '75_y', '76_y', '77_y', '78_y', '79_y',
      '80_y', '81_y', '82_y', '83_y', '84_y', '85_y', '86_y', '87_y',
      '88_y', '89_y', '90_y', '91_y', '92_y', '93_y', '94_y', '95_y'],
      dtype=object)
```

In [129]:

```
train_data_tfidf_w2v_tr = data[[ 'freq_qid1', 'freq_qid2',
    'q1len', 'q2len', 'q1_n_words', 'q2_n_words', 'word_Common',
    'word_Total', 'word_share', 'freq_q1+q2', 'freq_q1-q2', 'cwc_min',
    'cwc_max', 'csc_min', 'csc_max', 'ctc_min', 'ctc_max',
    'last_word_eq', 'first_word_eq', 'abs_len_diff', 'mean_len',
    'token_set_ratio', 'token_sort_ratio', 'fuzz_ratio',
    'fuzz_partial_ratio', 'longest_substr_ratio',
    '0_x', '1_x', '2_x', '3_x', '4_x', '5_x', '6_x',
    '7_x', '8_x', '9_x', '10_x', '11_x', '12_x', '13_x', '14_x',
    '15_x', '16_x', '17_x', '18_x', '19_x', '20_x', '21_x', '22_x',
    '23_x', '24_x', '25_x', '26_x', '27_x', '28_x', '29_x', '30_x',
    '31_x', '32_x', '33_x', '34_x', '35_x', '36_x', '37_x', '38_x',
    '39_x', '40_x', '41_x', '42_x', '43_x', '44_x', '45_x', '46_x',
    '47_x', '48_x', '49_x', '50_x', '51_x', '52_x', '53_x', '54_x',
    '55_x', '56_x', '57_x', '58_x', '59_x', '60_x', '61_x', '62_x',
    '63_x', '64_x', '65_x', '66_x', '67_x', '68_x', '69_x', '70_x',
    '71_x', '72_x', '73_x', '74_x', '75_x', '76_x', '77_x', '78_x',
    '79_x', '80_x', '81_x', '82_x', '83_x', '84_x', '85_x', '86_x',
    '87_x', '88_x', '89_x', '90_x', '91_x', '92_x', '93_x', '94_x',
    '95_x', '0_y', '1_y', '2_y', '3_y', '4_y', '5_y', '6_y', '7_y',
    '8_y', '9_y', '10_y', '11_y', '12_y', '13_y', '14_y', '15_y',
    '16_y', '17_y', '18_y', '19_y', '20_y', '21_y', '22_y', '23_y',
    '24_y', '25_y', '26_y', '27_y', '28_y', '29_y', '30_y', '31_y',
    '32_y', '33_y', '34_y', '35_y', '36_y', '37_y', '38_y', '39_y',
    '40_y', '41_y', '42_y', '43_y', '44_y', '45_y', '46_y', '47_y',
    '48_y', '49_y', '50_y', '51_y', '52_y', '53_y', '54_y', '55_y',
    '56_y', '57_y', '58_y', '59_y', '60_y', '61_y', '62_y', '63_y',
    '64_y', '65_y', '66_y', '67_y', '68_y', '69_y', '70_y', '71_y',
    '72_y', '73_y', '74_y', '75_y', '76_y', '77_y', '78_y', '79_y',
    '80_y', '81_y', '82_y', '83_y', '84_y', '85_y', '86_y', '87_y',
    '88_y', '89_y', '90_y', '91_y', '92_y', '93_y', '94_y', '95_y']]
```

In [130]:

```
train_data_tfidf_w2v_tr.head(2)
```

Out[130]:

	freq_qid1	freq_qid2	q1len	q2len	q1_n_words	q2_n_words	word_Common	word_Total	word_share	freq_q1+q2	freq
0	1	1	58	47	11	9	4.0	20.0	0.2	2	0
1	1	1	55	56	11	10	5.0	20.0	0.25	2	0

2 rows × 12 columns



In [131]:

```
train_data_tfidf_w2v_tr.shape, y_train.shape
```

Out[131]:

```
((7000, 218), (7000,))
```

## 8.1.2 Test data

In [132]:

```
#Creating db file from csv
import sqlalchemy
if not os.path.isfile(file_path + 'tfidf_w2v_test.db'):
    print("not present in drive")
    disk_engine = sqlalchemy.create_engine('sqlite:///'+file_path+'tfidf_w2v_test.db')
    start = dt.datetime.now()
    chunksize = 180000
    j = 0
```

```

index_start = 1
for df in pd.read_csv(file_path+'final_features_tfidf_w2v_test.csv',
                      names=[
                        'id', 'qid1', 'qid2', 'question1', 'question2', 'is_duplicate',
                        'freq_qid1', 'freq_qid2', 'q1len', 'q2len', 'q1_n_words',
                        'q2_n_words', 'word_Common', 'word_Total', 'word_share',
                        'freq_q1+q2', 'freq_q1-q2', 'cwc_min', 'cwc_max', 'csc_min',
                        'csc_max', 'ctc_min', 'ctc_max', 'last_word_eq', 'first_word_eq',
                        'abs_len_diff', 'mean_len', 'token_set_ratio', 'token_sort_ratio',
                        'fuzz_ratio', 'fuzz_partial_ratio', 'longest_substr_ratio',
                        'tfidf_w2v_q1', 'tfidf_w2v_q2', '0_x', '1_x', '2_x', '3_x', '4_x',
                        '5_x', '6_x', '7_x', '8_x', '9_x', '10_x', '11_x', '12_x', '13_x',
                        '14_x', '15_x', '16_x', '17_x', '18_x', '19_x', '20_x', '21_x',
                        '22_x', '23_x', '24_x', '25_x', '26_x', '27_x', '28_x', '29_x',
                        '30_x', '31_x', '32_x', '33_x', '34_x', '35_x', '36_x', '37_x',
                        '38_x', '39_x', '40_x', '41_x', '42_x', '43_x', '44_x', '45_x',
                        '46_x', '47_x', '48_x', '49_x', '50_x', '51_x', '52_x', '53_x',
                        '54_x', '55_x', '56_x', '57_x', '58_x', '59_x', '60_x', '61_x',
                        '62_x', '63_x', '64_x', '65_x', '66_x', '67_x', '68_x', '69_x',
                        '70_x', '71_x', '72_x', '73_x', '74_x', '75_x', '76_x', '77_x',
                        '78_x', '79_x', '80_x', '81_x', '82_x', '83_x', '84_x', '85_x',
                        '86_x', '87_x', '88_x', '89_x', '90_x', '91_x', '92_x', '93_x',
                        '94_x', '95_x', '0_y', '1_y', '2_y', '3_y', '4_y', '5_y', '6_y',
                        '7_y', '8_y', '9_y', '10_y', '11_y', '12_y', '13_y', '14_y',
                        '15_y', '16_y', '17_y', '18_y', '19_y', '20_y', '21_y', '22_y',
                        '23_y', '24_y', '25_y', '26_y', '27_y', '28_y', '29_y', '30_y',
                        '31_y', '32_y', '33_y', '34_y', '35_y', '36_y', '37_y', '38_y',
                        '39_y', '40_y', '41_y', '42_y', '43_y', '44_y', '45_y', '46_y',
                        '47_y', '48_y', '49_y', '50_y', '51_y', '52_y', '53_y', '54_y',
                        '55_y', '56_y', '57_y', '58_y', '59_y', '60_y', '61_y', '62_y',
                        '63_y', '64_y', '65_y', '66_y', '67_y', '68_y', '69_y', '70_y',
                        '71_y', '72_y', '73_y', '74_y', '75_y', '76_y', '77_y', '78_y',
                        '79_y', '80_y', '81_y', '82_y', '83_y', '84_y', '85_y', '86_y',
                        '87_y', '88_y', '89_y', '90_y', '91_y', '92_y', '93_y', '94_y',
                        '95_y'],
                      chunksize=chunksize, iterator=True, encoding='utf-8', ):
    df.index += index_start
    j+=1
    print('{} rows'.format(j*chunksize))
    df.to_sql('data', disk_engine, if_exists='append')
    index_start = df.index[-1] + 1

```

not present in drive  
180000 rows

In [133]:

```

read_db = file_path + 'tfidf_w2v_test.db'
conn_r = create_connection(read_db)
checkTableExists(conn_r)
conn_r.close()

```

Tables in the database:

```

data
[('data',)]

```

In [134]:

```

# try to sample data according to the computing power you have
if os.path.isfile(read_db):
    conn_r = create_connection(read_db)
    if conn_r is not None:
        # for selecting first 1M rows
        # data = pd.read_sql_query("""SELECT * FROM data LIMIT 100001;""", conn_r)
        data = pd.read_sql_query("""SELECT * From data ORDER BY RANDOM() LIMIT 100001;""", conn_r)

        # for selecting random points
        # data = pd.read_sql_query("SELECT * From data ORDER BY RANDOM() LIMIT 100001;", conn_r)
        conn_r.commit()
        conn_r.close()

```

In [135]:

```

data.shape

```



Out[135]:  
(3001, 227)

In [136]:  
data[data['is\_duplicate']=='is\_duplicate']

Out[136]:

	index	id	qid1	qid2	question1	question2	is_duplicate	freq_qid1	freq_qid2	q1len	q2len	q1_n_words	q2_n_words
1779	NaN	id	qid1	qid2	question1	question2	is_duplicate	freq_qid1	freq_qid2	q1len	q2len	q1_n_words	q2_n_words

1 rows × 227 columns

In [137]:  
data = data.drop(data.index[data['is\_duplicate']=='is\_duplicate'])

In [138]:  
data.shape

Out[138]:  
(3000, 227)

In [139]:  
`# remove the first row`  
`# data.drop(data.index[0], inplace=True)`  
`y_test = data['is_duplicate'].values`  
`data.drop(['id','index','is_duplicate'], axis=1, inplace=True)`

In [140]:  
data.shape, y\_test.shape

Out[140]:  
((3000, 224), (3000,))

In [141]:  
data.head(2)

Out[141]:

	qid1	qid2	question1	question2	freq_qid1	freq_qid2	q1len	q2len	q1_n_words	q2_n_words	word_Common	wo
0	480698	480699	should a skinny guy do cycling and exercise w...	do skinny people get drunk more easily	1	1	87	39	18	7	4.0	22.0
1	107697	107698	how did aristotle own and calilen	how do the beliefs of the atomic	1	1	78	77	12	13	4.0	24.0

	qid1	qid2	own question1 theories...	theory question2 from d...	freq_qid1	freq_qid2	q1len	q2len	q1_n_words	q2_n_words	word_Common	wo

2 rows × 224 columns

◀		▶
---	--	---

In [142]:

```
data.columns.values
```

Out[142]:

```
array(['qid1', 'qid2', 'question1', 'question2', 'freq_qid1', 'freq_qid2',
      'q1len', 'q2len', 'q1_n_words', 'q2_n_words', 'word_Common',
      'word_Total', 'word_share', 'freq_q1+q2', 'freq_q1-q2', 'cwc_min',
      'cwc_max', 'csc_min', 'csc_max', 'ctc_min', 'ctc_max',
      'last_word_eq', 'first_word_eq', 'abs_len_diff', 'mean_len',
      'token_set_ratio', 'token_sort_ratio', 'fuzz_ratio',
      'fuzz_partial_ratio', 'longest_substr_ratio', 'tfidf_w2v_q1',
      'tfidf_w2v_q2', '0_x', '1_x', '2_x', '3_x', '4_x', '5_x', '6_x',
      '7_x', '8_x', '9_x', '10_x', '11_x', '12_x', '13_x', '14_x',
      '15_x', '16_x', '17_x', '18_x', '19_x', '20_x', '21_x', '22_x',
      '23_x', '24_x', '25_x', '26_x', '27_x', '28_x', '29_x', '30_x',
      '31_x', '32_x', '33_x', '34_x', '35_x', '36_x', '37_x', '38_x',
      '39_x', '40_x', '41_x', '42_x', '43_x', '44_x', '45_x', '46_x',
      '47_x', '48_x', '49_x', '50_x', '51_x', '52_x', '53_x', '54_x',
      '55_x', '56_x', '57_x', '58_x', '59_x', '60_x', '61_x', '62_x',
      '63_x', '64_x', '65_x', '66_x', '67_x', '68_x', '69_x', '70_x',
      '71_x', '72_x', '73_x', '74_x', '75_x', '76_x', '77_x', '78_x',
      '79_x', '80_x', '81_x', '82_x', '83_x', '84_x', '85_x', '86_x',
      '87_x', '88_x', '89_x', '90_x', '91_x', '92_x', '93_x', '94_x',
      '95_x', '0_y', '1_y', '2_y', '3_y', '4_y', '5_y', '6_y', '7_y',
      '8_y', '9_y', '10_y', '11_y', '12_y', '13_y', '14_y', '15_y',
      '16_y', '17_y', '18_y', '19_y', '20_y', '21_y', '22_y', '23_y',
      '24_y', '25_y', '26_y', '27_y', '28_y', '29_y', '30_y', '31_y',
      '32_y', '33_y', '34_y', '35_y', '36_y', '37_y', '38_y', '39_y',
      '40_y', '41_y', '42_y', '43_y', '44_y', '45_y', '46_y', '47_y',
      '48_y', '49_y', '50_y', '51_y', '52_y', '53_y', '54_y', '55_y',
      '56_y', '57_y', '58_y', '59_y', '60_y', '61_y', '62_y', '63_y',
      '64_y', '65_y', '66_y', '67_y', '68_y', '69_y', '70_y', '71_y',
      '72_y', '73_y', '74_y', '75_y', '76_y', '77_y', '78_y', '79_y',
      '80_y', '81_y', '82_y', '83_y', '84_y', '85_y', '86_y', '87_y',
      '88_y', '89_y', '90_y', '91_y', '92_y', '93_y', '94_y', '95_y'],
      dtype=object)
```

In [143]:

```
test_data_tfidf_w2v_test = data[['freq_qid1', 'freq_qid2',
      'q1len', 'q2len', 'q1_n_words', 'q2_n_words', 'word_Common',
      'word_Total', 'word_share', 'freq_q1+q2', 'freq_q1-q2', 'cwc_min',
      'cwc_max', 'csc_min', 'csc_max', 'ctc_min', 'ctc_max',
      'last_word_eq', 'first_word_eq', 'abs_len_diff', 'mean_len',
      'token_set_ratio', 'token_sort_ratio', 'fuzz_ratio',
      'fuzz_partial_ratio', 'longest_substr_ratio',
      '0_x', '1_x', '2_x', '3_x', '4_x', '5_x', '6_x',
      '7_x', '8_x', '9_x', '10_x', '11_x', '12_x', '13_x', '14_x',
      '15_x', '16_x', '17_x', '18_x', '19_x', '20_x', '21_x', '22_x',
      '23_x', '24_x', '25_x', '26_x', '27_x', '28_x', '29_x', '30_x',
      '31_x', '32_x', '33_x', '34_x', '35_x', '36_x', '37_x', '38_x',
      '39_x', '40_x', '41_x', '42_x', '43_x', '44_x', '45_x', '46_x',
      '47_x', '48_x', '49_x', '50_x', '51_x', '52_x', '53_x', '54_x',
      '55_x', '56_x', '57_x', '58_x', '59_x', '60_x', '61_x', '62_x',
      '63_x', '64_x', '65_x', '66_x', '67_x', '68_x', '69_x', '70_x',
      '71_x', '72_x', '73_x', '74_x', '75_x', '76_x', '77_x', '78_x',
      '79_x', '80_x', '81_x', '82_x', '83_x', '84_x', '85_x', '86_x',
      '87_x', '88_x', '89_x', '90_x', '91_x', '92_x', '93_x', '94_x',
      '95_x', '0_y', '1_y', '2_y', '3_y', '4_y', '5_y', '6_y', '7_y',
      '8_y', '9_y', '10_y', '11_y', '12_y', '13_y', '14_y', '15_y',
      '16_y', '17_y', '18_y', '19_y', '20_y', '21_y', '22_y', '23_y',
      '24_y', '25_y', '26_y', '27_y', '28_y', '29_y', '30_y', '31_y',
      '32_y', '33_y', '34_y', '35_y', '36_y', '37_y', '38_y', '39_y',
      '40_y', '41_y', '42_y', '43_y', '44_y', '45_y', '46_y', '47_y',
      '48_y', '49_y', '50_y', '51_y', '52_y', '53_y', '54_y', '55_y',
      '56_y', '57_y', '58_y', '59_y', '60_y', '61_y', '62_y', '63_y',
      '64_y', '65_y', '66_y', '67_y', '68_y', '69_y', '70_y', '71_y',
      '72_y', '73_y', '74_y', '75_y', '76_y', '77_y', '78_y', '79_y',
      '80_y', '81_y', '82_y', '83_y', '84_y', '85_y', '86_y', '87_y',
      '88_y', '89_y', '90_y', '91_y', '92_y', '93_y', '94_y', '95_y']]
```

```
'64_y', '65_y', '66_y', '67_y', '68_y', '69_y', '70_y', '71_y',  
'72_y', '73_y', '74_y', '75_y', '76_y', '77_y', '78_y', '79_y',  
'80_y', '81_y', '82_y', '83_y', '84_y', '85_y', '86_y', '87_y',  
'88_y', '89_y', '90_y', '91_y', '92_y', '93_y', '94_y', '95_y']]
```

In [144]:

```
test_data_tfidf_w2v_test.head(2)
```

Out[144]:

	freq_qid1	freq_qid2	q1len	q2len	q1_n_words	q2_n_words	word_Common	word_Total	word_share	freq_q'
0	1	1	87	39	18	7	4.0	22.0	0.18181818181818182	2
1	1	1	78	77	12	13	4.0	24.0	0.16666666666666666	2

2 rows × 218 columns



In [145]:

```
test_data_tfidf_w2v_test.shape, y_test.shape
```

Out[145]:

```
((3000, 218), (3000,))
```

### 8.1.3 Converting String to Numerics

In [146]:

```
# after we read from sql table each entry was read it as a string  
# we convert all the features into numeric before we apply any model  
# https://stackoverflow.com/questions/40790031/pandas-to-numeric-find-out-which-string-it-was-unable-to-parse  
cols = list(train_data_tfidf_w2v_tr.columns)  
for i in cols:  
    train_data_tfidf_w2v_tr[i] = train_data_tfidf_w2v_tr[i].apply(pd.to_numeric, errors='coerce')  
    print(i)
```

```
freq_qid1  
freq_qid2  
q1len  
q2len  
q1_n_words  
q2_n_words  
word_Common  
word_Total  
word_share  
freq_q1+q2  
freq_q1-q2  
cwc_min  
cwc_max  
csc_min  
csc_max  
ctc_min  
ctc_max  
last_word_eq  
first_word_eq  
abs_len_diff  
mean_len  
token_set_ratio  
token_sort_ratio  
fuzz_ratio  
fuzz_partial_ratio  
longest_substr_ratio  
0_x  
1_v
```

1\_^  
2\_x  
3\_x  
4\_x  
5\_x  
6\_x  
7\_x  
8\_x  
9\_x  
10\_x  
11\_x  
12\_x  
13\_x  
14\_x  
15\_x  
16\_x  
17\_x  
18\_x  
19\_x  
20\_x  
21\_x  
22\_x  
23\_x  
24\_x  
25\_x  
26\_x  
27\_x  
28\_x  
29\_x  
30\_x  
31\_x  
32\_x  
33\_x  
34\_x  
35\_x  
36\_x  
37\_x  
38\_x  
39\_x  
40\_x  
41\_x  
42\_x  
43\_x  
44\_x  
45\_x  
46\_x  
47\_x  
48\_x  
49\_x  
50\_x  
51\_x  
52\_x  
53\_x  
54\_x  
55\_x  
56\_x  
57\_x  
58\_x  
59\_x  
60\_x  
61\_x  
62\_x  
63\_x  
64\_x  
65\_x  
66\_x  
67\_x  
68\_x  
69\_x  
70\_x  
71\_x  
72\_x  
73\_x  
74\_x  
75\_x  
76\_x  
77\_x  
78\_..

10\_x  
79\_x  
80\_x  
81\_x  
82\_x  
83\_x  
84\_x  
85\_x  
86\_x  
87\_x  
88\_x  
89\_x  
90\_x  
91\_x  
92\_x  
93\_x  
94\_x  
95\_x  
0\_y  
1\_y  
2\_y  
3\_y  
4\_y  
5\_y  
6\_y  
7\_y  
8\_y  
9\_y  
10\_y  
11\_y  
12\_y  
13\_y  
14\_y  
15\_y  
16\_y  
17\_y  
18\_y  
19\_y  
20\_y  
21\_y  
22\_y  
23\_y  
24\_y  
25\_y  
26\_y  
27\_y  
28\_y  
29\_y  
30\_y  
31\_y  
32\_y  
33\_y  
34\_y  
35\_y  
36\_y  
37\_y  
38\_y  
39\_y  
40\_y  
41\_y  
42\_y  
43\_y  
44\_y  
45\_y  
46\_y  
47\_y  
48\_y  
49\_y  
50\_y  
51\_y  
52\_y  
53\_y  
54\_y  
55\_y  
56\_y  
57\_y  
58\_y  
59\_y

```
59_y
60_y
61_y
62_y
63_y
64_y
65_y
66_y
67_y
68_y
69_y
70_y
71_y
72_y
73_y
74_y
75_y
76_y
77_y
78_y
79_y
80_y
81_y
82_y
83_y
84_y
85_y
86_y
87_y
88_y
89_y
90_y
91_y
92_y
93_y
94_y
95_y
```

In [147]:

```
y_train[:10]
```

Out[147]:

```
array(['1', '1', '1', '0', '0', '0', '1', '0', '1', '0'], dtype=object)
```

In [148]:

```
y_train = pd.Series(map(int,list(y_train)))
```

In [149]:

```
cols = list(test_data_tfidf_w2v_test.columns)
for i in cols:
    test_data_tfidf_w2v_test[i] = test_data_tfidf_w2v_test[i].apply(pd.to_numeric , errors='coerce'
)
    print(i)
```

```
freq_qid1
freq_qid2
qlen
q2len
q1_n_words
q2_n_words
word_Common
word_Total
word_share
freq_q1+q2
freq_q1-q2
cwc_min
cwc_max
csc_min
csc_max
etc_min
```

ccc\_max  
ctc\_max  
last\_word\_eq  
first\_word\_eq  
abs\_len\_diff  
mean\_len  
token\_set\_ratio  
token\_sort\_ratio  
fuzz\_ratio  
fuzz\_partial\_ratio  
longest\_substr\_ratio  
0\_x  
1\_x  
2\_x  
3\_x  
4\_x  
5\_x  
6\_x  
7\_x  
8\_x  
9\_x  
10\_x  
11\_x  
12\_x  
13\_x  
14\_x  
15\_x  
16\_x  
17\_x  
18\_x  
19\_x  
20\_x  
21\_x  
22\_x  
23\_x  
24\_x  
25\_x  
26\_x  
27\_x  
28\_x  
29\_x  
30\_x  
31\_x  
32\_x  
33\_x  
34\_x  
35\_x  
36\_x  
37\_x  
38\_x  
39\_x  
40\_x  
41\_x  
42\_x  
43\_x  
44\_x  
45\_x  
46\_x  
47\_x  
48\_x  
49\_x  
50\_x  
51\_x  
52\_x  
53\_x  
54\_x  
55\_x  
56\_x  
57\_x  
58\_x  
59\_x  
60\_x  
61\_x  
62\_x  
63\_x  
64\_x  
65\_x  
66\_x

67\_x  
68\_x  
69\_x  
70\_x  
71\_x  
72\_x  
73\_x  
74\_x  
75\_x  
76\_x  
77\_x  
78\_x  
79\_x  
80\_x  
81\_x  
82\_x  
83\_x  
84\_x  
85\_x  
86\_x  
87\_x  
88\_x  
89\_x  
90\_x  
91\_x  
92\_x  
93\_x  
94\_x  
95\_x  
0\_y  
1\_y  
2\_y  
3\_y  
4\_y  
5\_y  
6\_y  
7\_y  
8\_y  
9\_y  
10\_y  
11\_y  
12\_y  
13\_y  
14\_y  
15\_y  
16\_y  
17\_y  
18\_y  
19\_y  
20\_y  
21\_y  
22\_y  
23\_y  
24\_y  
25\_y  
26\_y  
27\_y  
28\_y  
29\_y  
30\_y  
31\_y  
32\_y  
33\_y  
34\_y  
35\_y  
36\_y  
37\_y  
38\_y  
39\_y  
40\_y  
41\_y  
42\_y  
43\_y  
44\_y  
45\_y  
46\_y  
47\_v



```
47_y
48_y
49_y
50_y
51_y
52_y
53_y
54_y
55_y
56_y
57_y
58_y
59_y
60_y
61_y
62_y
63_y
64_y
65_y
66_y
67_y
68_y
69_y
70_y
71_y
72_y
73_y
74_y
75_y
76_y
77_y
78_y
79_y
80_y
81_y
82_y
83_y
84_y
85_y
86_y
87_y
88_y
89_y
90_y
91_y
92_y
93_y
94_y
95_y
```

In [150]:

```
y_test[1:10]
```

Out[150]:

```
array(['0', '0', '0', '0', '0', '1', '1', '0', '0'], dtype=object)
```

In [151]:

```
y_test = pd.Series(map(int, list(y_test)))
```

In [152]:

```
x_tr_set1 = train_data_tfidf_w2v_tr
x_test_set1 = test_data_tfidf_w2v_test
y_tr_set1 = y_train
y_test_set1 = y_test
```

In [153]:

```
x_tr_set1.shape, x_test_set1.shape, y_tr_set1.shape, y_test_set1.shape
```

Out[153]:

```
((7000, 218), (3000, 218), (7000,), (3000,))
```

## 8.2 TFIDF

### 8.2.1 Train data

In [154]:

```
#Creating db file from csv
import sqlalchemy
if not os.path.isfile(file_path + 'tfidf_tr.db'):
    print("not present in drive")
    disk_engine = sqlalchemy.create_engine('sqlite:///'+file_path+'tfidf_tr.db')
    start = dt.datetime.now()
    chunksize = 180000
    j = 0
    index_start = 1
    for df in pd.read_csv(file_path+'final_features_tfidf_tr.csv',
                          names=['id', 'qid1', 'qid2', 'question1', 'question2', 'is_duplicate',
                                'freq_qid1', 'freq_qid2', 'qlen', 'q2len', 'q1_n_words',
                                'q2_n_words', 'word_Common', 'word_Total', 'word_share',
                                'freq_q1+q2', 'freq_q1-q2', 'cwc_min', 'cwc_max', 'csc_min',
                                'csc_max', 'ctc_min', 'ctc_max', 'last_word_eq', 'first_word_eq',
                                'abs_len_diff', 'mean_len', 'token_set_ratio', 'token_sort_ratio',
                                'fuzz_ratio', 'fuzz_partial_ratio', 'longest_substr_ratio',
                                '0_x', '1_x', '2_x', '3_x', '4_x',
                                '5_x', '6_x', '7_x', '8_x', '9_x', '10_x', '11_x', '12_x', '13_x',
                                '14_x', '15_x', '16_x', '17_x', '18_x', '19_x', '20_x', '21_x',
                                '22_x', '23_x', '24_x', '25_x', '26_x', '27_x', '28_x', '29_x',
                                '30_x', '31_x', '32_x', '33_x', '34_x', '35_x', '36_x', '37_x',
                                '38_x', '39_x', '40_x', '41_x', '42_x', '43_x', '44_x', '45_x',
                                '46_x', '47_x', '48_x', '49_x', '50_x', '51_x', '52_x', '53_x',
                                '54_x', '55_x', '56_x', '57_x', '58_x', '59_x', '60_x', '61_x',
                                '62_x', '63_x', '64_x', '65_x', '66_x', '67_x', '68_x', '69_x',
                                '70_x', '71_x', '72_x', '73_x', '74_x', '75_x', '76_x', '77_x',
                                '78_x', '79_x', '80_x', '81_x', '82_x', '83_x', '84_x', '85_x',
                                '86_x', '87_x', '88_x', '89_x', '90_x', '91_x', '92_x', '93_x',
                                '94_x', '95_x', '96_x', '97_x', '98_x', '99_x',
                                '0_y', '1_y', '2_y', '3_y', '4_y', '5_y', '6_y',
                                '7_y', '8_y', '9_y', '10_y', '11_y', '12_y', '13_y', '14_y',
                                '15_y', '16_y', '17_y', '18_y', '19_y', '20_y', '21_y', '22_y',
                                '23_y', '24_y', '25_y', '26_y', '27_y', '28_y', '29_y', '30_y',
                                '31_y', '32_y', '33_y', '34_y', '35_y', '36_y', '37_y', '38_y',
                                '39_y', '40_y', '41_y', '42_y', '43_y', '44_y', '45_y', '46_y',
                                '47_y', '48_y', '49_y', '50_y', '51_y', '52_y', '53_y', '54_y',
                                '55_y', '56_y', '57_y', '58_y', '59_y', '60_y', '61_y', '62_y',
                                '63_y', '64_y', '65_y', '66_y', '67_y', '68_y', '69_y', '70_y',
                                '71_y', '72_y', '73_y', '74_y', '75_y', '76_y', '77_y', '78_y',
                                '79_y', '80_y', '81_y', '82_y', '83_y', '84_y', '85_y', '86_y',
                                '87_y', '88_y', '89_y', '90_y', '91_y', '92_y', '93_y', '94_y',
                                '95_y', '96_y', '97_y', '98_y', '99_y'],
                          chunksize=chunksize, iterator=True, encoding='utf-8', ):
        df.index += index_start
        j+=1
        print('{} rows'.format(j*chunksize))
        df.to_sql('data', disk_engine, if_exists='append')
        index_start = df.index[-1] + 1
```

not present in drive  
180000 rows

In [155]:

```
read_db = file_path + 'tfidf_tr.db'
conn_r = create_connection(read_db)
checkTableExists(conn_r)
conn_r.close()
```

Tables in the database:

```
data
[('data',)]
```

In [156]:

```
# try to sample data according to the computing power you have
if os.path.isfile(read_db):
    conn_r = create_connection(read_db)
    if conn_r is not None:
        # for selecting first 1M rows
        # data = pd.read_sql_query("""SELECT * FROM data LIMIT 100001;""", conn_r)
        data = pd.read_sql_query("""SELECT * From data ORDER BY RANDOM() LIMIT 100001;""", conn_r)

        # for selecting random points
        # data = pd.read_sql_query("SELECT * From data ORDER BY RANDOM() LIMIT 100001;", conn_r)
    conn_r.commit()
    conn_r.close()
```

In [157]:

```
data.shape
```

Out[157]:

```
(7001, 233)
```

In [158]:

```
data[data['is_duplicate']=='is_duplicate']
```

Out[158]:

	index	id	qid1	qid2	question1	question2	is_duplicate	freq_qid1	freq_qid2	q1len	q2len	q1_n_words	q2_n_words
5848	NaN	id	qid1	qid2	question1	question2	is_duplicate	freq_qid1	freq_qid2	q1len	q2len	q1_n_words	q2_n_words

1 rows × 233 columns

In [159]:

```
data.head(2)
```

Out[159]:

	index	id	qid1	qid2	question1	question2	is_duplicate	freq_qid1	freq_qid2	q1len	q2len	q1_n_words	q2_n
0	711.0	244984	290121	9053	what are your favorite quotes or sayings	what are some of your favorite quotes	1	1	1	41	38	7	7
1	6942.0	230421	213704	32541	what is the difference between infatuation l...	how do you discern between infatuation and love	1	1	1	50	48	8	8

2 rows × 233 columns

In [160]:

```
data = data.drop(data.index[data['is_duplicate']=='is_duplicate'])
```

```
In [161]:
```

```
data.shape
```

```
Out[161]:
```

```
(7000, 233)
```

```
In [162]:
```

```
# remove the first row
# data.drop(data.index[0], inplace=True)
y_train = data['is_duplicate'].values
data.drop(['id', 'index', 'is_duplicate'], axis=1, inplace=True)
```

```
In [163]:
```

```
data.shape, y_train.shape
```

```
Out[163]:
```

```
((7000, 230), (7000,))
```

```
In [164]:
```

```
data.head(2)
```

```
Out[164]:
```

	qid1	qid2	question1	question2	freq_qid1	freq_qid2	q1len	q2len	q1_n_words	q2_n_words	word_Common	word_Total
0	290121	9053	what are your favorite quotes or sayings	what are some of your favorite quotes	1	1	41	38	7	7	4.0	14.0
1	213704	32541	what is the difference between infatuation l...	how do you discern between infatuation and love	1	1	50	48	8	8	3.0	16.0

2 rows × 230 columns



```
In [165]:
```

```
data.columns.values
```

```
Out[165]:
```

```
array(['qid1', 'qid2', 'question1', 'question2', 'freq_qid1', 'freq_qid2',
      'q1len', 'q2len', 'q1_n_words', 'q2_n_words', 'word_Common',
      'word_Total', 'word_share', 'freq_q1+q2', 'freq_q1-q2', 'cwc_min',
      'cwc_max', 'csc_min', 'csc_max', 'ctc_min', 'ctc_max',
      'last_word_eq', 'first_word_eq', 'abs_len_diff', 'mean_len',
      'token_set_ratio', 'token_sort_ratio', 'fuzz_ratio',
      'fuzz_partial_ratio', 'longest_substr_ratio', '0_x', '1_x', '2_x',
      '3_x', '4_x', '5_x', '6_x', '7_x', '8_x', '9_x', '10_x', '11_x',
      '12_x', '13_x', '14_x', '15_x', '16_x', '17_x', '18_x', '19_x',
      '20_x', '21_x', '22_x', '23_x', '24_x', '25_x', '26_x', '27_x',
      '28_x', '29_x', '30_x', '31_x', '32_x', '33_x', '34_x', '35_x',
      '36_x', '37_x', '38_x', '39_x', '40_x', '41_x', '42_x', '43_x',
      '44_x', '45_x', '46_x', '47_x', '48_x', '49_x', '50_x', '51_x',
      '52_x', '53_x', '54_x', '55_x', '56_x', '57_x', '58_x', '59_x',
      '60_x', '61_x', '62_x', '63_x', '64_x', '65_x', '66_x', '67_x',
```

```
'68_x', '69_x', '70_x', '71_x', '72_x', '73_x', '74_x', '75_x',
'76_x', '77_x', '78_x', '79_x', '80_x', '81_x', '82_x', '83_x',
'84_x', '85_x', '86_x', '87_x', '88_x', '89_x', '90_x', '91_x',
'92_x', '93_x', '94_x', '95_x', '96_x', '97_x', '98_x', '99_x',
'0_y', '1_y', '2_y', '3_y', '4_y', '5_y', '6_y', '7_y', '8_y',
'9_y', '10_y', '11_y', '12_y', '13_y', '14_y', '15_y', '16_y',
'17_y', '18_y', '19_y', '20_y', '21_y', '22_y', '23_y', '24_y',
'25_y', '26_y', '27_y', '28_y', '29_y', '30_y', '31_y', '32_y',
'33_y', '34_y', '35_y', '36_y', '37_y', '38_y', '39_y', '40_y',
'41_y', '42_y', '43_y', '44_y', '45_y', '46_y', '47_y', '48_y',
'49_y', '50_y', '51_y', '52_y', '53_y', '54_y', '55_y', '56_y',
'57_y', '58_y', '59_y', '60_y', '61_y', '62_y', '63_y', '64_y',
'65_y', '66_y', '67_y', '68_y', '69_y', '70_y', '71_y', '72_y',
'73_y', '74_y', '75_y', '76_y', '77_y', '78_y', '79_y', '80_y',
'81_y', '82_y', '83_y', '84_y', '85_y', '86_y', '87_y', '88_y',
'89_y', '90_y', '91_y', '92_y', '93_y', '94_y', '95_y', '96_y',
'97_y', '98_y', '99_y'], dtype=object)
```

In [166]:

```
train_data_tfidf_tr = data[[ 'freq_qid1', 'freq_qid2',
'q1len', 'q2len', 'q1_n_words', 'q2_n_words', 'word_Common',
'word_Total', 'word_share', 'freq_q1+q2', 'freq_q1-q2', 'cwc_min',
'cwc_max', 'csc_min', 'csc_max', 'ctc_min', 'ctc_max',
'last_word_eq', 'first_word_eq', 'abs_len_diff', 'mean_len',
'token_set_ratio', 'token_sort_ratio', 'fuzz_ratio',
'fuzz_partial_ratio', 'longest_substr_ratio', '0_x', '1_x', '2_x',
'3_x', '4_x', '5_x', '6_x', '7_x', '8_x', '9_x', '10_x', '11_x',
'12_x', '13_x', '14_x', '15_x', '16_x', '17_x', '18_x', '19_x',
'20_x', '21_x', '22_x', '23_x', '24_x', '25_x', '26_x', '27_x',
'28_x', '29_x', '30_x', '31_x', '32_x', '33_x', '34_x', '35_x',
'36_x', '37_x', '38_x', '39_x', '40_x', '41_x', '42_x', '43_x',
'44_x', '45_x', '46_x', '47_x', '48_x', '49_x', '50_x', '51_x',
'52_x', '53_x', '54_x', '55_x', '56_x', '57_x', '58_x', '59_x',
'60_x', '61_x', '62_x', '63_x', '64_x', '65_x', '66_x', '67_x',
'68_x', '69_x', '70_x', '71_x', '72_x', '73_x', '74_x', '75_x',
'76_x', '77_x', '78_x', '79_x', '80_x', '81_x', '82_x', '83_x',
'84_x', '85_x', '86_x', '87_x', '88_x', '89_x', '90_x', '91_x',
'92_x', '93_x', '94_x', '95_x', '96_x', '97_x', '98_x', '99_x',
'0_y', '1_y', '2_y', '3_y', '4_y', '5_y', '6_y', '7_y', '8_y',
'9_y', '10_y', '11_y', '12_y', '13_y', '14_y', '15_y', '16_y',
'17_y', '18_y', '19_y', '20_y', '21_y', '22_y', '23_y', '24_y',
'25_y', '26_y', '27_y', '28_y', '29_y', '30_y', '31_y', '32_y',
'33_y', '34_y', '35_y', '36_y', '37_y', '38_y', '39_y', '40_y',
'41_y', '42_y', '43_y', '44_y', '45_y', '46_y', '47_y', '48_y',
'49_y', '50_y', '51_y', '52_y', '53_y', '54_y', '55_y', '56_y',
'57_y', '58_y', '59_y', '60_y', '61_y', '62_y', '63_y', '64_y',
'65_y', '66_y', '67_y', '68_y', '69_y', '70_y', '71_y', '72_y',
'73_y', '74_y', '75_y', '76_y', '77_y', '78_y', '79_y', '80_y',
'81_y', '82_y', '83_y', '84_y', '85_y', '86_y', '87_y', '88_y',
'89_y', '90_y', '91_y', '92_y', '93_y', '94_y', '95_y', '96_y',
'97_y', '98_y', '99_y']]
```

In [167]:

```
train_data_tfidf_tr.head(2)
```

Out[167]:

	freq_qid1	freq_qid2	q1len	q2len	q1_n_words	q2_n_words	word_Common	word_Total	word_share	freq_q1-
0	1	1	41	38	7	7	4.0	14.0	0.2857142857142857	2
1	1	1	50	48	8	8	3.0	16.0	0.1875	2

2 rows × 226 columns



In [168]:

```
train_data_tfidf_tr.shape, y_train.shape
```

.....

Out[168]:

```
((7000, 226), (7000,))
```

## 8.2.2 Test data

In [169]:

```
#Creating db file from csv
import sqlalchemy
if not os.path.isfile(file_path + 'tfidf_test.db'):
    print("not present in drive")
    disk_engine = sqlalchemy.create_engine('sqlite:///'+file_path+'tfidf_test.db')
    start = dt.datetime.now()
    chunksize = 180000
    j = 0
    index_start = 1
    for df in pd.read_csv(file_path+'final_features_tfidf_test.csv',
                          names=['id', 'qid1', 'qid2', 'question1', 'question2', 'is_duplicate',
                                'freq_qid1', 'freq_qid2', 'qlen', 'q2len', 'q1_n_words',
                                'q2_n_words', 'word_Common', 'word_Total', 'word_share',
                                'freq_q1+q2', 'freq_q1-q2', 'cwc_min', 'cwc_max', 'csc_min',
                                'csc_max', 'ctc_min', 'ctc_max', 'last_word_eq', 'first_word_eq',
                                'abs_len_diff', 'mean_len', 'token_set_ratio', 'token_sort_ratio',
                                'fuzz_ratio', 'fuzz_partial_ratio', 'longest_substr_ratio',
                                '0_x', '1_x', '2_x', '3_x', '4_x',
                                '5_x', '6_x', '7_x', '8_x', '9_x', '10_x', '11_x', '12_x', '13_x',
                                '14_x', '15_x', '16_x', '17_x', '18_x', '19_x', '20_x', '21_x',
                                '22_x', '23_x', '24_x', '25_x', '26_x', '27_x', '28_x', '29_x',
                                '30_x', '31_x', '32_x', '33_x', '34_x', '35_x', '36_x', '37_x',
                                '38_x', '39_x', '40_x', '41_x', '42_x', '43_x', '44_x', '45_x',
                                '46_x', '47_x', '48_x', '49_x', '50_x', '51_x', '52_x', '53_x',
                                '54_x', '55_x', '56_x', '57_x', '58_x', '59_x', '60_x', '61_x',
                                '62_x', '63_x', '64_x', '65_x', '66_x', '67_x', '68_x', '69_x',
                                '70_x', '71_x', '72_x', '73_x', '74_x', '75_x', '76_x', '77_x',
                                '78_x', '79_x', '80_x', '81_x', '82_x', '83_x', '84_x', '85_x',
                                '86_x', '87_x', '88_x', '89_x', '90_x', '91_x', '92_x', '93_x',
                                '94_x', '95_x', '96_x', '97_x', '98_x', '99_x',
                                '0_y', '1_y', '2_y', '3_y', '4_y', '5_y', '6_y',
                                '7_y', '8_y', '9_y', '10_y', '11_y', '12_y', '13_y', '14_y',
                                '15_y', '16_y', '17_y', '18_y', '19_y', '20_y', '21_y', '22_y',
                                '23_y', '24_y', '25_y', '26_y', '27_y', '28_y', '29_y', '30_y',
                                '31_y', '32_y', '33_y', '34_y', '35_y', '36_y', '37_y', '38_y',
                                '39_y', '40_y', '41_y', '42_y', '43_y', '44_y', '45_y', '46_y',
                                '47_y', '48_y', '49_y', '50_y', '51_y', '52_y', '53_y', '54_y',
                                '55_y', '56_y', '57_y', '58_y', '59_y', '60_y', '61_y', '62_y',
                                '63_y', '64_y', '65_y', '66_y', '67_y', '68_y', '69_y', '70_y',
                                '71_y', '72_y', '73_y', '74_y', '75_y', '76_y', '77_y', '78_y',
                                '79_y', '80_y', '81_y', '82_y', '83_y', '84_y', '85_y', '86_y',
                                '87_y', '88_y', '89_y', '90_y', '91_y', '92_y', '93_y', '94_y',
                                '95_y', '96_y', '97_y', '98_y', '99_y'],
                          chunksize=chunksize, iterator=True, encoding='utf-8', ):
        df.index += index_start
        j+=1
        print('{} rows'.format(j*chunksize))
        df.to_sql('data', disk_engine, if_exists='append')
        index_start = df.index[-1] + 1
```

not present in drive  
180000 rows

In [170]:

```
read_db = file_path + 'tfidf_test.db'
conn_r = create_connection(read_db)
checkTableExists(conn_r)
conn_r.close()
```

Tables in the database:  
data  
[('data',)]

In [171]:

```
# try to sample data according to the computing power you have
if os.path.isfile(read_db):
    conn_r = create_connection(read_db)
    if conn_r is not None:
        # for selecting first 1M rows
        # data = pd.read_sql_query("""SELECT * FROM data LIMIT 100001;""", conn_r)
        data =pd.read_sql_query("""SELECT * From data ORDER BY RANDOM() LIMIT 100001;""", conn_r)

        # for selecting random points
        # data = pd.read_sql_query("SELECT * From data ORDER BY RANDOM() LIMIT 100001;", conn_r)
    conn_r.commit()
    conn_r.close()
```

In [172]:

```
data.shape
```

Out[172]:

```
(3001, 233)
```

In [173]:

```
data[data['is_duplicate']=='is_duplicate']
```

Out[173]:

	index	id	qid1	qid2	question1	question2	is_duplicate	freq_qid1	freq_qid2	q1len	q2len	q1_n_words	q2_n_words
2443	NaN	id	qid1	qid2	question1	question2	is_duplicate	freq_qid1	freq_qid2	q1len	q2len	q1_n_words	q2_n_words

1 rows × 233 columns



In [174]:

```
data = data.drop(data.index[data['is_duplicate']=='is_duplicate'])
```

In [175]:

```
data.shape
```

Out[175]:

```
(3000, 233)
```

In [176]:

```
# remove the first row
# data.drop(data.index[0], inplace=True)
y_test = data['is_duplicate'].values
data.drop(['id','index','is_duplicate'], axis=1, inplace=True)
```

In [177]:

```
data.shape, y_test.shape
```

Out[177]:

```
((3000, 230), (3000,))
```

In [178]:

```
data.head(2)
```

Out [178]:

	qid1	qid2	question1	question2	freq_qid1	freq_qid2	q1len	q2len	q1_n_words	q2_n_words	word_Common	wo
0	84277	123591	why even though my nose is oily the skin kee...	what causes skin to peel off	1	1	86	29	17	6	3.0	23.0
1	394136	394137	what are the best law colleges in india	which is the best college for law in india	1	1	40	43	8	9	5.0	17.0

2 rows x 230 columns

In [179]:

```
data.columns.values
```

Out [179] :

```
array(['qid1', 'qid2', 'question1', 'question2', 'freq_qid1', 'freq_qid2',
      'q1len', 'q2len', 'q1_n_words', 'q2_n_words', 'word_Common',
      'word_Total', 'word_share', 'freq_q1+q2', 'freq_q1-q2', 'cwc_min',
      'cwc_max', 'csc_min', 'csc_max', 'ctc_min', 'ctc_max',
      'last_word_eq', 'first_word_eq', 'abs_len_diff', 'mean_len',
      'token_set_ratio', 'token_sort_ratio', 'fuzz_ratio',
      'fuzz_partial_ratio', 'longest_substr_ratio', '0_x', '1_x', '2_x',
      '3_x', '4_x', '5_x', '6_x', '7_x', '8_x', '9_x', '10_x', '11_x',
      '12_x', '13_x', '14_x', '15_x', '16_x', '17_x', '18_x', '19_x',
      '20_x', '21_x', '22_x', '23_x', '24_x', '25_x', '26_x', '27_x',
      '28_x', '29_x', '30_x', '31_x', '32_x', '33_x', '34_x', '35_x',
      '36_x', '37_x', '38_x', '39_x', '40_x', '41_x', '42_x', '43_x',
      '44_x', '45_x', '46_x', '47_x', '48_x', '49_x', '50_x', '51_x',
      '52_x', '53_x', '54_x', '55_x', '56_x', '57_x', '58_x', '59_x',
      '60_x', '61_x', '62_x', '63_x', '64_x', '65_x', '66_x', '67_x',
      '68_x', '69_x', '70_x', '71_x', '72_x', '73_x', '74_x', '75_x',
      '76_x', '77_x', '78_x', '79_x', '80_x', '81_x', '82_x', '83_x',
      '84_x', '85_x', '86_x', '87_x', '88_x', '89_x', '90_x', '91_x',
      '92_x', '93_x', '94_x', '95_x', '96_x', '97_x', '98_x', '99_x',
      '0_y', '1_y', '2_y', '3_y', '4_y', '5_y', '6_y', '7_y', '8_y',
      '9_y', '10_y', '11_y', '12_y', '13_y', '14_y', '15_y', '16_y',
      '17_y', '18_y', '19_y', '20_y', '21_y', '22_y', '23_y', '24_y',
      '25_y', '26_y', '27_y', '28_y', '29_y', '30_y', '31_y', '32_y',
      '33_y', '34_y', '35_y', '36_y', '37_y', '38_y', '39_y', '40_y',
      '41_y', '42_y', '43_y', '44_y', '45_y', '46_y', '47_y', '48_y',
      '49_y', '50_y', '51_y', '52_y', '53_y', '54_y', '55_y', '56_y',
      '57_y', '58_y', '59_y', '60_y', '61_y', '62_y', '63_y', '64_y',
      '65_y', '66_y', '67_y', '68_y', '69_y', '70_y', '71_y', '72_y',
      '73_y', '74_y', '75_y', '76_y', '77_y', '78_y', '79_y', '80_y',
      '81_y', '82_y', '83_y', '84_y', '85_y', '86_y', '87_y', '88_y',
      '89_y', '90_y', '91_y', '92_y', '93_y', '94_y', '95_y', '96_y',
      '97_y', '98_y', '99_y'], dtype=object)
```

In [180]:

```
test_data_tfidf_test = data[[ 'freq_qid1', 'freq_qid2',
                              'q1len', 'q2len', 'q1_n_words', 'q2_n_words', 'word_Common',
                              'word_Total', 'word_share', 'freq_q1+q2', 'freq_q1-q2', 'cwc_min',
                              'cwc_max', 'csc_min', 'csc_max', 'ctc_min', 'ctc_max',
                              'last_word_eq', 'first_word_eq', 'abs_len_diff', 'mean_len',
                              'token_set_ratio', 'token_sort_ratio', 'fuzz_ratio',
                              'fuzz_partial_ratio', 'longest_substr_ratio', '0_x', '1_x', '2_x',
                              '3_x', '4_x', '5_x', '6_x', '7_x', '8_x', '9_x', '10_x', '11_x',
                              '12_x', '13_x', '14_x', '15_x', '16_x', '17_x', '18_x', '19_x',
                              '20_x', '21_x', '22_x', '23_x', '24_x', '25_x', '26_x', '27_x',
                              '28_x', '29_x', '30_x', '31_x', '32_x', '33_x', '34_x', '35_x',
                              '36_x', '37_x', '38_x', '39_x', '40_x', '41_x', '42_x', '43_x',
```



```
'44_x', '45_x', '46_x', '47_x', '48_x', '49_x', '50_x', '51_x',
'52_x', '53_x', '54_x', '55_x', '56_x', '57_x', '58_x', '59_x',
'60_x', '61_x', '62_x', '63_x', '64_x', '65_x', '66_x', '67_x',
'68_x', '69_x', '70_x', '71_x', '72_x', '73_x', '74_x', '75_x',
'76_x', '77_x', '78_x', '79_x', '80_x', '81_x', '82_x', '83_x',
'84_x', '85_x', '86_x', '87_x', '88_x', '89_x', '90_x', '91_x',
'92_x', '93_x', '94_x', '95_x', '96_x', '97_x', '98_x', '99_x',
'0_y', '1_y', '2_y', '3_y', '4_y', '5_y', '6_y', '7_y', '8_y',
'9_y', '10_y', '11_y', '12_y', '13_y', '14_y', '15_y', '16_y',
'17_y', '18_y', '19_y', '20_y', '21_y', '22_y', '23_y', '24_y',
'25_y', '26_y', '27_y', '28_y', '29_y', '30_y', '31_y', '32_y',
'33_y', '34_y', '35_y', '36_y', '37_y', '38_y', '39_y', '40_y',
'41_y', '42_y', '43_y', '44_y', '45_y', '46_y', '47_y', '48_y',
'49_y', '50_y', '51_y', '52_y', '53_y', '54_y', '55_y', '56_y',
'57_y', '58_y', '59_y', '60_y', '61_y', '62_y', '63_y', '64_y',
'65_y', '66_y', '67_y', '68_y', '69_y', '70_y', '71_y', '72_y',
'73_y', '74_y', '75_y', '76_y', '77_y', '78_y', '79_y', '80_y',
'81_y', '82_y', '83_y', '84_y', '85_y', '86_y', '87_y', '88_y',
'89_y', '90_y', '91_y', '92_y', '93_y', '94_y', '95_y', '96_y',
'97_y', '98_y', '99_y']]
```

In [181]:

```
test_data_tfidf_test.head(2)
```

Out[181]:

	freq_qid1	freq_qid2	q1len	q2len	q1_n_words	q2_n_words	word_Common	word_Total	word_share	freq_q'
0	1	1	86	29	17	6	3.0	23.0	0.13043478260869565	2
1	1	1	40	43	8	9	5.0	17.0	0.29411764705882354	2

2 rows × 226 columns



In [182]:

```
test_data_tfidf_test.shape, y_test.shape
```

Out[182]:

```
((3000, 226), (3000,))
```

## 8.2.3 Converting String to Numerics

In [183]:

```
# after we read from sql table each entry was read it as a string
# we convert all the features into numeric before we apply any model
# https://stackoverflow.com/questions/40790031/pandas-to-numeric-find-out-which-string-it-was-unable-to-parse
cols = list(train_data_tfidf_tr.columns)
for i in cols:
    train_data_tfidf_tr[i] = train_data_tfidf_tr[i].apply(pd.to_numeric, errors='coerce')
    print(i)
```

```
freq_qid1
freq_qid2
q1len
q2len
q1_n_words
q2_n_words
word_Common
word_Total
word_share
freq_q1+q2
freq_q1-q2
cwc_min
cwc_max
csc min
```

```
...-
csc_max
ctc_min
ctc_max
last_word_eq
first_word_eq
abs_len_diff
mean_len
token_set_ratio
token_sort_ratio
fuzz_ratio
fuzz_partial_ratio
longest_substr_ratio
0_x
1_x
2_x
3_x
4_x
5_x
6_x
7_x
8_x
9_x
10_x
11_x
12_x
13_x
14_x
15_x
16_x
17_x
18_x
19_x
20_x
21_x
22_x
23_x
24_x
25_x
26_x
27_x
28_x
29_x
30_x
31_x
32_x
33_x
34_x
35_x
36_x
37_x
38_x
39_x
40_x
41_x
42_x
43_x
44_x
45_x
46_x
47_x
48_x
49_x
50_x
51_x
52_x
53_x
54_x
55_x
56_x
57_x
58_x
59_x
60_x
61_x
62_x
63_x
64_x
```

~!-..  
65\_x  
66\_x  
67\_x  
68\_x  
69\_x  
70\_x  
71\_x  
72\_x  
73\_x  
74\_x  
75\_x  
76\_x  
77\_x  
78\_x  
79\_x  
80\_x  
81\_x  
82\_x  
83\_x  
84\_x  
85\_x  
86\_x  
87\_x  
88\_x  
89\_x  
90\_x  
91\_x  
92\_x  
93\_x  
94\_x  
95\_x  
96\_x  
97\_x  
98\_x  
99\_x  
0\_y  
1\_y  
2\_y  
3\_y  
4\_y  
5\_y  
6\_y  
7\_y  
8\_y  
9\_y  
10\_y  
11\_y  
12\_y  
13\_y  
14\_y  
15\_y  
16\_y  
17\_y  
18\_y  
19\_y  
20\_y  
21\_y  
22\_y  
23\_y  
24\_y  
25\_y  
26\_y  
27\_y  
28\_y  
29\_y  
30\_y  
31\_y  
32\_y  
33\_y  
34\_y  
35\_y  
36\_y  
37\_y  
38\_y  
39\_y  
40\_y  
41\_v

```
42_y
43_y
44_y
45_y
46_y
47_y
48_y
49_y
50_y
51_y
52_y
53_y
54_y
55_y
56_y
57_y
58_y
59_y
60_y
61_y
62_y
63_y
64_y
65_y
66_y
67_y
68_y
69_y
70_y
71_y
72_y
73_y
74_y
75_y
76_y
77_y
78_y
79_y
80_y
81_y
82_y
83_y
84_y
85_y
86_y
87_y
88_y
89_y
90_y
91_y
92_y
93_y
94_y
95_y
96_y
97_y
98_y
99_y
```

In [184]:

```
y_train[:10]
```

Out[184]:

```
array(['1', '1', '0', '1', '0', '0', '0', '0', '0', '0'], dtype=object)
```

In [185]:

```
y_train = pd.Series(map(int,list(y_train)))
```

In [186]:

```
cols = list(test_data.tfidf_test.columns)
```

```
cols = list(test_data_tfidf_test.columns)
for i in cols:
    test_data_tfidf_test[i] = test_data_tfidf_test[i].apply(pd.to_numeric , errors='coerce')
print(i)
```

```
freq_qid1
freq_qid2
q1len
q2len
q1_n_words
q2_n_words
word_Common
word_Total
word_share
freq_q1+q2
freq_q1-q2
cwc_min
cwc_max
csc_min
csc_max
ctc_min
ctc_max
last_word_eq
first_word_eq
abs_len_diff
mean_len
token_set_ratio
token_sort_ratio
fuzz_ratio
fuzz_partial_ratio
longest_substr_ratio
0_x
1_x
2_x
3_x
4_x
5_x
6_x
7_x
8_x
9_x
10_x
11_x
12_x
13_x
14_x
15_x
16_x
17_x
18_x
19_x
20_x
21_x
22_x
23_x
24_x
25_x
26_x
27_x
28_x
29_x
30_x
31_x
32_x
33_x
34_x
35_x
36_x
37_x
38_x
39_x
40_x
41_x
42_x
43_x
44_x
45_x
```

46\_x  
47\_x  
48\_x  
49\_x  
50\_x  
51\_x  
52\_x  
53\_x  
54\_x  
55\_x  
56\_x  
57\_x  
58\_x  
59\_x  
60\_x  
61\_x  
62\_x  
63\_x  
64\_x  
65\_x  
66\_x  
67\_x  
68\_x  
69\_x  
70\_x  
71\_x  
72\_x  
73\_x  
74\_x  
75\_x  
76\_x  
77\_x  
78\_x  
79\_x  
80\_x  
81\_x  
82\_x  
83\_x  
84\_x  
85\_x  
86\_x  
87\_x  
88\_x  
89\_x  
90\_x  
91\_x  
92\_x  
93\_x  
94\_x  
95\_x  
96\_x  
97\_x  
98\_x  
99\_x  
0\_y  
1\_y  
2\_y  
3\_y  
4\_y  
5\_y  
6\_y  
7\_y  
8\_y  
9\_y  
10\_y  
11\_y  
12\_y  
13\_y  
14\_y  
15\_y  
16\_y  
17\_y  
18\_y  
19\_y  
20\_y  
21\_y  
22\_y

23\_y  
24\_y  
25\_y  
26\_y  
27\_y  
28\_y  
29\_y  
30\_y  
31\_y  
32\_y  
33\_y  
34\_y  
35\_y  
36\_y  
37\_y  
38\_y  
39\_y  
40\_y  
41\_y  
42\_y  
43\_y  
44\_y  
45\_y  
46\_y  
47\_y  
48\_y  
49\_y  
50\_y  
51\_y  
52\_y  
53\_y  
54\_y  
55\_y  
56\_y  
57\_y  
58\_y  
59\_y  
60\_y  
61\_y  
62\_y  
63\_y  
64\_y  
65\_y  
66\_y  
67\_y  
68\_y  
69\_y  
70\_y  
71\_y  
72\_y  
73\_y  
74\_y  
75\_y  
76\_y  
77\_y  
78\_y  
79\_y  
80\_y  
81\_y  
82\_y  
83\_y  
84\_y  
85\_y  
86\_y  
87\_y  
88\_y  
89\_y  
90\_y  
91\_y  
92\_y  
93\_y  
94\_y  
95\_y  
96\_y  
97\_y  
98\_y  
99\_v

In [187]:

```
y_test[1:10]
```

Out[187]:

```
array(['1', '0', '0', '1', '1', '1', '1', '0', '0'], dtype=object)
```

In [188]:

```
y_test = pd.Series(map(int, list(y_test)))
```

In [189]:

```
x_tr_set2 = train_data_tfidf_tr
x_test_set2 = test_data_tfidf_test
y_tr_set2 = y_train
y_test_set2 = y_test
```

In [190]:

```
y_test
```

Out[190]:

```
0      0
1      1
2      0
3      0
4      1
..
2995   0
2996   1
2997   0
2998   1
2999   0
Length: 3000, dtype: int64
```

In [191]:

```
x_tr_set1.shape, x_test_set1.shape, y_tr_set1.shape, y_test_set1.shape
```

Out[191]:

```
((7000, 218), (3000, 218), (7000,), (3000,))
```

In [192]:

```
x_tr_set2.shape, x_test_set2.shape, y_tr_set2.shape, y_test_set2.shape
```

Out[192]:

```
((7000, 226), (3000, 226), (7000,), (3000,))
```

## 9.Models

In [193]:

```
# This function plots the confusion matrices given y_i, y_i_hat.
def plot_confusion_matrix(test_y, predict_y):
    C = confusion_matrix(test_y, predict_y)
    # C = 9,9 matrix, each cell (i,j) represents number of points of class i are predicted class j

    A = ((C.T) / (C.sum(axis=1))).T
    #divid each element of the confusion matrix with the sum of elements in that column

    # C = [[1  2]
```



```

# C = [[1, 2],
#       [3, 4]]
# C.T = [[1, 3],
#         [2, 4]]
# C.sum(axis = 1)  axis=0 corresponds to columns and axis=1 corresponds to rows in two
dimensional array
# C.sum(axis = 1) = [[3, 7]]
# ((C.T)/(C.sum(axis=1))) = [[1/3, 3/7],
#                             [2/3, 4/7]]

# ((C.T)/(C.sum(axis=1))).T = [[1/3, 2/3],
#                               [3/7, 4/7]]
# sum of row elements = 1

B = (C/C.sum(axis=0))
#divid each element of the confusion matrix with the sum of elements in that row
# C = [[1, 2],
#       [3, 4]]
# C.sum(axis = 0)  axis=0 corresponds to columns and axis=1 corresponds to rows in two
dimensional array
# C.sum(axis = 0) = [[4, 6]]
# (C/C.sum(axis=0)) = [[1/4, 2/6],
#                       [3/4, 4/6]]
plt.figure(figsize=(20,4))

labels = [1,2]
# representing A in heatmap format
cmap=sns.light_palette("blue")
plt.subplot(1, 3, 1)
sns.heatmap(C, annot=True, cmap=cmap, fmt=".3f", xticklabels=labels, yticklabels=labels)
plt.xlabel('Predicted Class')
plt.ylabel('Original Class')
plt.title("Confusion matrix")

plt.subplot(1, 3, 2)
sns.heatmap(B, annot=True, cmap=cmap, fmt=".3f", xticklabels=labels, yticklabels=labels)
plt.xlabel('Predicted Class')
plt.ylabel('Original Class')
plt.title("Precision matrix")

plt.subplot(1, 3, 3)
# representing B in heatmap format
sns.heatmap(A, annot=True, cmap=cmap, fmt=".3f", xticklabels=labels, yticklabels=labels)
plt.xlabel('Predicted Class')
plt.ylabel('Original Class')
plt.title("Recall matrix")

plt.show()

```

## 9.1 Random model

### 9.1.1 on Set1 TFIDF-W2V

In [194]:

```

# we need to generate 9 numbers and the sum of numbers should be 1
# one solution is to generate 9 numbers and divide each of the numbers by their sum
# ref: https://stackoverflow.com/a/18662466/4084039
# we create a output array that has exactly same size as the CV data

train_len = len(y_tr_set1)
predicted_y = np.zeros((train_len,2))
for i in range(train_len):
    rand_probs = np.random.rand(1,2)
    predicted_y[i] = ((rand_probs/sum(sum(rand_probs))))[0])

loss_tr = log_loss(y_tr_set1, predicted_y, eps=1e-15)
print("Log loss on Train Data using Random Model",log_loss(y_tr_set1, predicted_y, eps=1e-15))

test_len = len(y_test_set1)
predicted_y = np.zeros((test_len,2))
for i in range(test_len):
    rand_probs = np.random.rand(1,2)
    predicted_y[i] = ((rand_probs/sum(sum(rand_probs))))[0])

```

```

predicted_y[i] = ((rand_probs/sum(sum(rand_probs)))[0])


loss_test = log_loss(y_test_set1, predicted_y, eps=1e-15)
print("Log loss on Test Data using Random Model",log_loss(y_test_set1, predicted_y, eps=1e-15))

predicted_y = np.argmax(predicted_y, axis=1)
plot_confusion_matrix(y_test, predicted_y)

```

Log loss on Train Data using Random Model 0.8909857413369713

Log loss on Test Data using Random Model 0.8839580037167362



In [195]:

```

from prettytable import PrettyTable
summary = PrettyTable()
summary.field_names = ["Vectorizer", "Model", "Best Hyperparameters", "Log loss on Train", "Log loss on Test"]
summary.add_row(['TFIDF_w2v','Random dumb model', "-", loss_tr, loss_test ])
print(summary)

```

Vectorizer	Model	Best Hyperparameters	Log loss on Train	Log loss on Test
TFIDF_w2v	Random dumb model	-	0.8909857413369713	0.8839580037167362

## 9.1.2 on Set2 TFIDF

In [196]:

```

# we need to generate 9 numbers and the sum of numbers should be 1
# one solution is to generate 9 numbers and divide each of the numbers by their sum
# ref: https://stackoverflow.com/a/18662466/4084039
# we create a output array that has exactly same size as the CV data

train_len = len(y_tr_set2)
predicted_y = np.zeros((train_len,2))
for i in range(train_len):
    rand_probs = np.random.rand(1,2)
    predicted_y[i] = ((rand_probs/sum(sum(rand_probs)))[0])

loss_tr = log_loss(y_tr_set2, predicted_y, eps=1e-15)
print("Log loss on Train Data using Random Model",log_loss(y_tr_set2, predicted_y, eps=1e-15))

test_len = len(y_test_set2)
predicted_y = np.zeros((test_len,2))
for i in range(test_len):
    rand_probs = np.random.rand(1,2)
    predicted_y[i] = ((rand_probs/sum(sum(rand_probs)))[0])


loss_test = log_loss(y_test_set2, predicted_y, eps=1e-15)
print("Log loss on Test Data using Random Model",log_loss(y_test_set2, predicted_y, eps=1e-15))

predicted_y = np.argmax(predicted_y, axis=1)
plot_confusion_matrix(y_test, predicted_y)

```

Log loss on Train Data using Random Model 0.8681386222435589

Log loss on Test Data using Random Model 0.9077642168454928



In [197]:

```
# from prettytable import PrettyTable
# summary = PrettyTable()
summary.field_names = ["Vectorizer", "Model", "Best Hyperparameters", "Log loss on Train", "Log loss on Test"]
summary.add_row(['TFIDF', 'Random dumb model', "-", loss_tr, loss_test ])
print(summary)
```

Vectorizer	Model	Best Hyperparameters	Log loss on Train	Log loss on Test
TFIDF_w2v	Random dumb model	-	0.8909857413369713	0.8839580037167362
TFIDF	Random dumb model	-	0.8681386222435589	0.9077642168454928

## 9.2 Logistic Regression with hyperparameter tuning

### 9.2.1 on Set1 TFIDF-W2V

In [198]:

```
alpha = [10 ** x for x in range(-5, 2)] # hyperparam for SGD classifier.

# read more about SGDClassifier() at http://scikit-learn.org/stable/modules/generated/sklearn.linear_model.SGDClassifier.html
# -----
# default parameters
# SGDClassifier(loss='hinge', penalty='l2', alpha=0.0001, l1_ratio=0.15, fit_intercept=True, max_iter=None, tol=None,
# shuffle=True, verbose=0, epsilon=0.1, n_jobs=1, random_state=None, learning_rate='optimal', eta0=0.0, power_t=0.5,
# class_weight=None, warm_start=False, average=False, n_iter=None)

# some of methods
# fit(X, y[, coef_init, intercept_init, ...]) Fit linear model with Stochastic Gradient Descent.
# predict(X) Predict class labels for samples in X.

#-----
# video link:
#-----

log_error_test = []
log_error_tr = []

for c,i in enumerate(alpha):
    clf = SGDClassifier(alpha=i, penalty='l2', loss='log', random_state=42)
    clf.fit(x_tr_set1, y_tr_set1)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(x_tr_set1, y_tr_set1)

    predict_y = sig_clf.predict_proba(x_tr_set1)
    log_error_tr.append(log_loss(y_tr_set1, predict_y, labels=clf.classes_, eps=1e-15))

    predict_y = sig_clf.predict_proba(x_test_set1)
    log_error_test.append(log_loss(y_test_set1, predict_y, labels=clf.classes_, eps=1e-15))
    print("="*40)
    print('values of alpha = ', i)
```

```

print(values_of_alpha[-1], 1)
print("The Train log loss is:", log_error_tr[c])
print("The Test log loss is:", log_error_test[c])

```

```

=====
values of alpha = 1e-05
The Train log loss is: 0.6091916455448889
The Test log loss is: 0.6176956398711984
=====
values of alpha = 0.0001
The Train log loss is: 0.543096536913335
The Test log loss is: 0.5677820606079191
=====
values of alpha = 0.001
The Train log loss is: 0.5633423389191436
The Test log loss is: 0.5794821211753041
=====
values of alpha = 0.01
The Train log loss is: 0.5394195102762249
The Test log loss is: 0.5632093221739694
=====
values of alpha = 0.1
The Train log loss is: 0.5181022850363174
The Test log loss is: 0.5504720745790216
=====
values of alpha = 1
The Train log loss is: 0.5332981263234295
The Test log loss is: 0.567115087691264
=====
values of alpha = 10
The Train log loss is: 0.5473516900016429
The Test log loss is: 0.5788983703201077

```


In [199]:

```

fig, ax = plt.subplots()
ax.plot(alpha, log_error_test, label="test")
# for i, txt in enumerate(np.round(log_error_array,3)):
#     ax.annotate((alpha[i],np.round(txt,3)), (alpha[i],log_error_array[i]))

ax.plot(alpha, log_error_tr, label="train")
# for i, txt in enumerate(np.round(log_error_array,3)):
#     ax.annotate((alpha[i],np.round(txt,3)), (alpha[i],log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.legend()
plt.show()

```



In [200]:

```

best_alpha = np.argmin(log_error_test)
alpha[best_alpha]

```

Out[200]:

0.1

In [201]:

```
clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l2', loss='log', random_state=42)
clf.fit(x_tr_set1, y_tr_set1)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(x_tr_set1, y_tr_set1)


predict_y = sig_clf.predict_proba(x_tr_set1)
loss_tr = log_loss(y_tr_set1, predict_y, labels=clf.classes_, eps=1e-15)
print('For values of best alpha = ', alpha[best_alpha], "The train log loss
is:", log_loss(y_tr_set1, predict_y, labels=clf.classes_, eps=1e-15))

predict_y = sig_clf.predict_proba(x_test_set1)
loss_test = log_loss(y_test_set1, predict_y, labels=clf.classes_, eps=1e-15)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss
is:", log_loss(y_test_set1, predict_y, labels=clf.classes_, eps=1e-15))

predicted_y = np.argmax(predict_y, axis=1)
print("Total number of data points :", len(predicted_y))

plot_confusion_matrix(y_test_set1, predicted_y)
```

For values of best alpha = 0.1 The train log loss is: 0.5181022850363174  
For values of best alpha = 0.1 The test log loss is: 0.5504720745790216  
Total number of data points : 3000



In [202]:

```
# summary.field_names = ["Vectorizer", "Model", "Best Hyperparameters", "Log loss on Train", "Log
loss on Test"]
summary.add_row(['TFIDF_w2v', 'Logistic Regression by iterative search', "alpha = " +
str(alpha[best_alpha]) , loss_tr, loss_test ])
print(summary)
```

Vectorizer	Model	Best Hyperparameters	Log loss on Train
TFIDF_w2v	Random dumb model	-	0.8909857413369713
TFIDF	Random dumb model	-	0.8681386222435589
TFIDF_w2v	Logistic Regression by iterative search	alpha = 0.1	0.5181022850363174

## 9.2.2 on Set2 TFIDF

In [203]:

```
alpha = [10 ** x for x in range(-5, 2)] # hyperparam for SGD classifier.

# read more about SGDClassifier() at http://scikit-
```

```

learn.org/stable/modules/generated/sklearn.linear_model.SGDClassifier.html
# -----
# default parameters
# SGDClassifier(loss='hinge', penalty='l2', alpha=0.0001, l1_ratio=0.15, fit_intercept=True, max_iter=None, tol=None,
# shuffle=True, verbose=0, epsilon=0.1, n_jobs=1, random_state=None, learning_rate='optimal', eta0=0.0, power_t=0.5,
# class_weight=None, warm_start=False, average=False, n_iter=None)

# some of methods
# fit(X, y[, coef_init, intercept_init, ...]) Fit linear model with Stochastic Gradient Descent.
# predict(X) Predict class labels for samples in X.

#-----
# video link:
#-----

log_error_test = []
log_error_tr = []

for c,i in enumerate(alpha):
    clf = SGDClassifier(alpha=i, penalty='l2', loss='log', random_state=42)
    clf.fit(x_tr_set2, y_tr_set2)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(x_tr_set2, y_tr_set2)

    predict_y = sig_clf.predict_proba(x_tr_set2)
    log_error_tr.append(log_loss(y_tr_set2, predict_y, labels=clf.classes_, eps=1e-15))

    predict_y = sig_clf.predict_proba(x_test_set2)
    log_error_test.append(log_loss(y_test_set2, predict_y, labels=clf.classes_, eps=1e-15))
    print("="*40)
    print('values of alpha = ', i)
    print( "The Train log loss is:", log_error_tr[c])
    print( "The Test log loss is:", log_error_test[c])

```

```

=====
values of alpha = 1e-05
The Train log loss is: 0.5033447327977969
The Test log loss is: 0.518974269660145
=====
values of alpha = 0.0001
The Train log loss is: 0.5136749062267901
The Test log loss is: 0.5246955453491132
=====
values of alpha = 0.001
The Train log loss is: 0.5004491052297945
The Test log loss is: 0.5149092374073378
=====
values of alpha = 0.01
The Train log loss is: 0.507871340132494
The Test log loss is: 0.5198158058284845
=====
values of alpha = 0.1
The Train log loss is: 0.538881494767776
The Test log loss is: 0.547810219482023
=====
values of alpha = 1
The Train log loss is: 0.5708097348310461
The Test log loss is: 0.576765884165689
=====
values of alpha = 10
The Train log loss is: 0.5917558991396166
The Test log loss is: 0.598314345312015

```

In [204]:


```

fig, ax = plt.subplots()
ax.plot(alpha, log_error_test, label="test")
# for i, txt in enumerate(np.round(log_error_array,3)):
#     ax.annotate((alpha[i],np.round(txt,3)), (alpha[i],log_error_array[i]))

ax.plot(alpha, log_error_tr, label="train")
# for i, txt in enumerate(np.round(log_error_array,3)):
#     ax.annotate((alpha[i],np.round(txt,3)), (alpha[i],log_error_array[i]))

```

```
ax.annotate([alpha[i], np.round(log_e, 3)], (alpha[i], log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.legend()
plt.show()
```



In [205]:

```
best_alpha = np.argmin(log_error_test)
alpha[best_alpha]
```

Out[205]:

0.001

In [206]:

```
clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l2', loss='log', random_state=42)
clf.fit(x_tr_set2, y_tr_set2)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(x_tr_set2, y_tr_set2)


predict_y = sig_clf.predict_proba(x_tr_set2)
loss_tr = log_loss(y_tr_set2, predict_y, labels=clf.classes_, eps=1e-15)
print('For values of best alpha = ', alpha[best_alpha], "The train log loss is:", log_loss(y_tr_set2, predict_y, labels=clf.classes_, eps=1e-15))

predict_y = sig_clf.predict_proba(x_test_set2)
loss_test = log_loss(y_test_set2, predict_y, labels=clf.classes_, eps=1e-15)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:", log_loss(y_test_set2, predict_y, labels=clf.classes_, eps=1e-15))

predicted_y = np.argmax(predict_y, axis=1)
print("Total number of data points :", len(predicted_y))

plot_confusion_matrix(y_test_set2, predicted_y)
```

For values of best alpha = 0.001 The train log loss is: 0.5004491052297945  
 For values of best alpha = 0.001 The test log loss is: 0.5149092374073378  
 Total number of data points : 3000



In [207]:

```
# summary.field_names = ["Vectorizer", "Model", "Best Hyperparameters", "Log loss on Train", "Log  
loss on Test"]  
summary.add_row(['TFIDF', 'Logistic Regression by iterative search', "alpha = " +  
str(alpha[best_alpha]) , loss_tr, loss_test ])  
print(summary)
```

Vectorizer	Model	Best Hyperparameters	Log loss on Train
Log loss on Test			
TFIDF_w2v	Random dumb model	-	0.8909857413369713
0.8839580037167362			
TFIDF	Random dumb model	-	0.8681386222435589
0.9077642168454928			
TFIDF_w2v	Logistic Regression by iterative search	alpha = 0.1	0.5181022850363174
0.5504720745790216			
TFIDF	Logistic Regression by iterative search	alpha = 0.001	0.5004491052297945
0.5149092374073378			

## 9.3 Linear SVM with hyperparameter tuning

### 9.3.1 on Set1 TFIDF-W2V

#### 9.3.1.1 using L1 regularizer

In [208]:

```
alpha = [10 ** x for x in range(-5, 2)] # hyperparam for SGD classifier.  
  
# read more about SGDClassifier() at http://scikit-learn.org/stable/modules/generated/sklearn.linear\_model.SGDClassifier.html  
# -----  
# default parameters  
# SGDClassifier(loss='hinge', penalty='l2', alpha=0.0001, l1_ratio=0.15, fit_intercept=True, max_iter=None, tol=None,  
# shuffle=True, verbose=0, epsilon=0.1, n_jobs=1, random_state=None, learning_rate='optimal', eta0=0.0, power_t=0.5,  
# class_weight=None, warm_start=False, average=False, n_iter=None)  
  
# some of methods  
# fit(X, y[, coef_init, intercept_init, ...]) Fit linear model with Stochastic Gradient Descent.  
# predict(X) Predict class labels for samples in X.  
  
#-----  
# video link:  
#-----  
  
log_error_test = []  
log_error_tr = []  
  
for c,i in enumerate(alpha):  
    clf = SGDClassifier(alpha=i, penalty='l1', loss='hinge', random_state=42)  
    clf.fit(x_tr_set1, y_tr_set1)  
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")  
    sig_clf.fit(x_tr_set1, y_tr_set1)  
  
    predict_y = sig_clf.predict_proba(x_tr_set1)  
    log_error_tr.append(log_loss(y_tr_set1, predict_y, labels=clf.classes_, eps=1e-15))  
  
    predict_y = sig_clf.predict_proba(x_test_set1)  
    log_error_test.append(log_loss(y_test_set1, predict_y, labels=clf.classes_, eps=1e-15))  
    print("="*40)  
    print('values of alpha = ', i)  
    print("The Train log loss is:", log_error_tr[c])
```




```
print( "The Test log loss is:", log_error_test[c])
```

```
=====
values of alpha = 1e-05
The Train log loss is: 0.6599367618737666
The Test log loss is: 0.660011544065058
=====
values of alpha = 0.0001
The Train log loss is: 0.527068330818554
The Test log loss is: 0.5528737395236669
=====
values of alpha = 0.001
The Train log loss is: 0.5312839392947547
The Test log loss is: 0.5620421712320395
=====
values of alpha = 0.01
The Train log loss is: 0.5331349687886794
The Test log loss is: 0.558945396873255
=====
values of alpha = 0.1
The Train log loss is: 0.5956886197710739
The Test log loss is: 0.6080006713561955
=====
values of alpha = 1
The Train log loss is: 0.6339287504368991
The Test log loss is: 0.6360698125216653
=====
values of alpha = 10
The Train log loss is: 0.659919827255779
The Test log loss is: 0.6597429897754892
```

In [209]:

```
fig, ax = plt.subplots()
ax.plot(alpha, log_error_test, label="test")
# for i, txt in enumerate(np.round(log_error_array,3)):
#     ax.annotate((alpha[i],np.round(txt,3)), (alpha[i],log_error_array[i]))

ax.plot(alpha, log_error_tr, label="train")
# for i, txt in enumerate(np.round(log_error_array,3)):
#     ax.annotate((alpha[i],np.round(txt,3)), (alpha[i],log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.legend()
plt.show()
```



In [210]:

```
best_alpha = np.argmin(log_error_test)
alpha[best_alpha]
```

Out[210]:

0.0001

In [211]:

```
clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l1', loss='hinge', random_state=42)
clf.fit(x_tr_set1, y_tr_set1)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(x_tr_set1, y_tr_set1)


predict_y = sig_clf.predict_proba(x_tr_set1)
loss_tr = log_loss(y_tr_set1, predict_y, labels=clf.classes_, eps=1e-15)
print('For values of best alpha = ', alpha[best_alpha], "The train log loss
is:", log_loss(y_tr_set1, predict_y, labels=clf.classes_, eps=1e-15))

predict_y = sig_clf.predict_proba(x_test_set1)
loss_test = log_loss(y_test_set1, predict_y, labels=clf.classes_, eps=1e-15)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss
is:", log_loss(y_test_set1, predict_y, labels=clf.classes_, eps=1e-15))

predicted_y = np.argmax(predict_y, axis=1)
print("Total number of data points :", len(predicted_y))

plot_confusion_matrix(y_test_set1, predicted_y)
```

For values of best alpha = 0.0001 The train log loss is: 0.527068330818554  
For values of best alpha = 0.0001 The test log loss is: 0.5528737395236669  
Total number of data points : 3000



In [212]:

```
# summary.field_names = ["Vectorizer", "Model", "Best Hyperparameters", "Log loss on Train", "Log
loss on Test"]
summary.add_row(['TFIDF_w2v', 'Linear SVM by iterative search', "L1 Norm & alpha = " + str(alpha[b
est_alpha]), loss_tr, loss_test])
print(summary)
```

Vectorizer	Model	Best Hyperparameters	Log loss on Train	Log loss on Test
TFIDF_w2v	Random dumb model	-	-	0.8909857413369713
TFIDF	Random dumb model	-	-	0.8839580037167362
TFIDF_w2v	Logistic Regression by iterative search	alpha = 0.1	-	0.9077642168454928
TFIDF	Logistic Regression by iterative search	alpha = 0.001	-	0.5181022850363174
TFIDF_w2v	Linear SVM by iterative search	L1 Norm & alpha = 0.0001	0.527068330818554	0.5504720745790216

### 9.3.1.2 using L2 regularizer

In [213]:

```
alpha = [10 ** x for x in range(-5, 2)] # hyperparam for SGD classifier.
```

```

# read more about SGDClassifier() at http://scikit-
learn.org/stable/modules/generated/sklearn.linear_model.SGDClassifier.html
# -----
# default parameters
# SGDClassifier(loss='hinge', penalty='l2', alpha=0.0001, l1_ratio=0.15, fit_intercept=True, max_i
ter=None, tol=None,
# shuffle=True, verbose=0, epsilon=0.1, n_jobs=1, random_state=None, learning_rate='optimal', eta0
=0.0, power_t=0.5,
# class_weight=None, warm_start=False, average=False, n_iter=None)

# some of methods
# fit(X, y[, coef_init, intercept_init, ...]) Fit linear model with Stochastic Gradient Descent.
# predict(X) Predict class labels for samples in X.

#-----
# video link:
#-----

```

```

log_error_test = []
log_error_tr = []

for c,i in enumerate(alpha):
    clf = SGDClassifier(alpha=i, penalty='l2', loss='hinge', random_state=42)
    clf.fit(x_tr_set1, y_tr_set1)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(x_tr_set1, y_tr_set1)

    predict_y = sig_clf.predict_proba(x_tr_set1)
    log_error_tr.append(log_loss(y_tr_set1, predict_y, labels=clf.classes_, eps=1e-15))

    predict_y = sig_clf.predict_proba(x_test_set1)
    log_error_test.append(log_loss(y_test_set1, predict_y, labels=clf.classes_, eps=1e-15))
    print(" "*40)
    print('values of alpha = ', i)
    print( "The Train log loss is:", log_error_tr[c])
    print( "The Test log loss is:", log_error_test[c])

```

```

=====
values of alpha = 1e-05
The Train log loss is: 0.6599367618737666
The Test log loss is: 0.660011544065058
=====
values of alpha = 0.0001
The Train log loss is: 0.5483462114713233
The Test log loss is: 0.5734025274118046
=====
values of alpha = 0.001
The Train log loss is: 0.5463185742896588
The Test log loss is: 0.5680875320477938
=====
values of alpha = 0.01
The Train log loss is: 0.5422897955597696
The Test log loss is: 0.5653708871224005
=====
values of alpha = 0.1
The Train log loss is: 0.5287501978376552
The Test log loss is: 0.5569509511703374
=====
values of alpha = 1
The Train log loss is: 0.5341690990494615
The Test log loss is: 0.5664107320888319
=====
values of alpha = 10
The Train log loss is: 0.5399361664670905
The Test log loss is: 0.5730685406280625

```

In [214]:

```


fig, ax = plt.subplots()
ax.plot(alpha, log_error_test, label="test")
# for i, txt in enumerate(np.round(log_error_array,3)):
#     ax.annotate((alpha[i],np.round(txt,3)), (alpha[i],log_error_array[i]))

```

```

ax.plot(alpha, log_error_tr, label="train")
# for i, txt in enumerate(np.round(log_error_array,3)):
#     ax.annotate((alpha[i],np.round(txt,3)), (alpha[i],log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.legend()
plt.show()

```



In [215]:

```

best_alpha = np.argmin(log_error_test)
alpha[best_alpha]

```

Out[215]:

0.1

In [216]:

```

clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l1', loss='hinge', random_state=42)
clf.fit(x_tr_set1, y_tr_set1)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(x_tr_set1, y_tr_set1)

predict_y = sig_clf.predict_proba(x_tr_set1)
loss_tr = log_loss(y_tr_set1, predict_y, labels=clf.classes_, eps=1e-15)
print('For values of best alpha = ', alpha[best_alpha], "The train log loss is:", log_loss(y_tr_set1, predict_y, labels=clf.classes_, eps=1e-15))


predict_y = sig_clf.predict_proba(x_test_set1)
loss_test = log_loss(y_test_set1, predict_y, labels=clf.classes_, eps=1e-15)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:", log_loss(y_test_set1, predict_y, labels=clf.classes_, eps=1e-15))

predicted_y = np.argmax(predict_y,axis=1)
print("Total number of data points :", len(predicted_y))

plot_confusion_matrix(y_test_set1, predicted_y)

```

For values of best alpha = 0.1 The train log loss is: 0.5956886197710739  
For values of best alpha = 0.1 The test log loss is: 0.6080006713561955  
Total number of data points : 3000



In [217]:

```
# summary.field_names = ["Vectorizer", "Model", "Best Hyperparameters", "Log loss on Train", "Log
loss on Test"]
summary.add_row(['TFIDF_w2v', 'Linear SVM by iterative search', "L2 Norm & alpha = " + str(alpha[best_alpha]), loss_tr, loss_test])
print(summary)
```

Vectorizer	Model	Best Hyperparameters	Log loss on Train	Log loss on Test
TFIDF_w2v	Random dumb model	-	0.8909857413369713	0.8839580037167362
TFIDF	Random dumb model	-	0.9077642168454928	0.86813862224359
TFIDF_w2v	Logistic Regression by iterative search	alpha = 0.1	0.5181022850363174	0.5504720745790216
TFIDF	Logistic Regression by iterative search	alpha = 0.001	0.5004491052297945	0.5149092374073378
TFIDF_w2v	Linear SVM by iterative search	L1 Norm & alpha = 0.0001	0.5528737395236669	0.52706833081854
TFIDF_w2v	Linear SVM by iterative search	L2 Norm & alpha = 0.1	0.6080006713561955	0.5956886197710739

## 9.3.2 on Set2 TFIDF

### 9.3.2.1 using L1 regularizer

In [218]:

```
alpha = [10 ** x for x in range(-5, 2)] # hyperparam for SGD classifier.

# read more about SGDClassifier() at http://scikit-learn.org/stable/modules/generated/sklearn.linear_model.SGDClassifier.html
# -----
# default parameters
# SGDClassifier(loss='hinge', penalty='l2', alpha=0.0001, l1_ratio=0.15, fit_intercept=True, max_iter=None, tol=None,
# shuffle=True, verbose=0, epsilon=0.1, n_jobs=1, random_state=None, learning_rate='optimal', eta0=0.0, power_t=0.5,
# class_weight=None, warm_start=False, average=False, n_iter=None)

# some of methods
# fit(X, y[, coef_init, intercept_init, ...]) Fit linear model with Stochastic Gradient Descent.
# predict(X) Predict class labels for samples in X.

#-----
# video link:
#-----

log_error_test = []
log_error_tr = []

for c,i in enumerate(alpha):
    clf = SGDClassifier(alpha=i, penalty='l1', loss='hinge', random_state=42)
    clf.fit(x_tr_set2, y_tr_set2)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(x_tr_set2, y_tr_set2)

    predict_y = sig_clf.predict_proba(x_tr_set2)
    log_error_tr.append(log_loss(y_tr_set2, predict_y, labels=clf.classes_, eps=1e-15))

    predict_y = sig_clf.predict_proba(x_test_set2)
    log_error_test.append(log_loss(y_test_set2, predict_y, labels=clf.classes_, eps=1e-15))
```

```

print("="*40)
print('values of alpha = ', i)
print( "The Train log loss is:", log_error_tr[c])
print( "The Test log loss is:", log_error_test[c])

```

```

=====
values of alpha = 1e-05
The Train log loss is: 0.49838297495456246
The Test log loss is: 0.5148041636287534
=====
values of alpha = 0.0001
The Train log loss is: 0.5304182143472248
The Test log loss is: 0.544454054175011
=====
values of alpha = 0.001
The Train log loss is: 0.5389807644172705
The Test log loss is: 0.551159201065387
=====
values of alpha = 0.01
The Train log loss is: 0.5418008693703278
The Test log loss is: 0.5509921907365074
=====
values of alpha = 0.1
The Train log loss is: 0.5699004693783398
The Test log loss is: 0.573834749185796
=====
values of alpha = 1
The Train log loss is: 0.6077947237022424
The Test log loss is: 0.6128287980616041
=====
values of alpha = 10
The Train log loss is: 0.6404419011995035
The Test log loss is: 0.6452703803045354

```


In [219]:

```

fig, ax = plt.subplots()
ax.plot(alpha, log_error_test, label="test")
# for i, txt in enumerate(np.round(log_error_array,3)):
#     ax.annotate((alpha[i],np.round(txt,3)), (alpha[i],log_error_array[i]))

ax.plot(alpha, log_error_tr, label="train")
# for i, txt in enumerate(np.round(log_error_array,3)):
#     ax.annotate((alpha[i],np.round(txt,3)), (alpha[i],log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.legend()
plt.show()

```



In [220]:

```

best_alpha = np.argmin(log_error_test)
alpha[best_alpha]

```

Out[220]:

1e-05

In [221]:

```
clf = SGDCClassifier(alpha=alpha[best_alpha], penalty='l1', loss='hinge', random_state=42)
clf.fit(x_tr_set2, y_tr_set2)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(x_tr_set2, y_tr_set2)

predict_y = sig_clf.predict_proba(x_tr_set2)
loss_tr = log_loss(y_tr_set2, predict_y, labels=clf.classes_, eps=1e-15)
print('For values of best alpha = ', alpha[best_alpha], "The train log loss
is:", log_loss(y_tr_set2, predict_y, labels=clf.classes_, eps=1e-15))

predict_y = sig_clf.predict_proba(x_test_set2)
loss_test = log_loss(y_test_set2, predict_y, labels=clf.classes_, eps=1e-15)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss
is:", log_loss(y_test_set2, predict_y, labels=clf.classes_, eps=1e-15))


predicted_y = np.argmax(predict_y, axis=1)
print("Total number of data points :", len(predicted_y))

plot_confusion_matrix(y_test_set2, predicted_y)
```

For values of best alpha = 1e-05 The train log loss is: 0.49838297495456246

For values of best alpha = 1e-05 The test log loss is: 0.5148041636287534

Total number of data points : 3000



In [222]:

```
# summary.field_names = ["Vectorizer", "Model", "Best Hyperparameters", "Log loss on Train", "Log
loss on Test"]
summary.add_row(['TFIDF', 'Linear SVM by iterative search', "L1 Norm & alpha = " +
str(alpha[best_alpha]) , loss_tr, loss_test ])
print(summary)
```

-----				
-+-----+				
Vectorizer		Model	Best Hyperparameters	Log loss on Tr
in   Log loss on Test				
+-----+				
-+-----+				
TFIDF_w2v		Random dumb model	-	
0.8909857413369713		0.8839580037167362		
TFIDF		Random dumb model	-	0.868138622243
89   0.9077642168454928				
TFIDF_w2v		Logistic Regression by iterative search	alpha = 0.1	
0.5181022850363174		0.5504720745790216		
TFIDF		Logistic Regression by iterative search	alpha = 0.001	
0.5004491052297945		0.5149092374073378		
TFIDF_w2v		Linear SVM by iterative search	L1 Norm & alpha = 0.0001	0.527068330818
554   0.5528737395236669				
TFIDF_w2v		Linear SVM by iterative search	L2 Norm & alpha = 0.1	0.595688619771
0739   0.6080006713561955				
TFIDF		Linear SVM by iterative search	L1 Norm & alpha = 1e-05	
0.49838297495456246		0.5148041636287534		
+-----+				
-+-----+				

### 9.3.2.2 using L2 regularizer

In [223]:

```
alpha = [10 ** x for x in range(-5, 2)] # hyperparam for SGD classifier.

# read more about SGDClassifier() at http://scikit-learn.org/stable/modules/generated/sklearn.linear_model.SGDClassifier.html
# -----
# default parameters
# SGDClassifier(loss='hinge', penalty='l2', alpha=0.0001, l1_ratio=0.15, fit_intercept=True, max_iter=None, tol=None,
# shuffle=True, verbose=0, epsilon=0.1, n_jobs=1, random_state=None, learning_rate='optimal', eta0=0.0, power_t=0.5,
# class_weight=None, warm_start=False, average=False, n_iter=None)

# some of methods
# fit(X, y[, coef_init, intercept_init, ...]) Fit linear model with Stochastic Gradient Descent.
# predict(X) Predict class labels for samples in X.

#-----
# video link:
#-----

log_error_test = []
log_error_tr = []

for c,i in enumerate(alpha):
    clf = SGDClassifier(alpha=i, penalty='l2', loss='hinge', random_state=42)
    clf.fit(x_tr_set2, y_tr_set2)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(x_tr_set2, y_tr_set2)

    predict_y = sig_clf.predict_proba(x_tr_set2)
    log_error_tr.append(log_loss(y_tr_set2, predict_y, labels=clf.classes_, eps=1e-15))

    predict_y = sig_clf.predict_proba(x_test_set2)
    log_error_test.append(log_loss(y_test_set2, predict_y, labels=clf.classes_, eps=1e-15))
    print("="*40)
    print('values of alpha = ', i)
    print("The Train log loss is:", log_error_tr[c])
    print("The Test log loss is:", log_error_test[c])


=====
values of alpha = 1e-05
The Train log loss is: 0.5047591525799148
The Test log loss is: 0.5180016919378869
=====
values of alpha = 0.0001
The Train log loss is: 0.5013833062480888
The Test log loss is: 0.5145545037573294
=====
values of alpha = 0.001
The Train log loss is: 0.5006735023574405
The Test log loss is: 0.51441888795156
=====
values of alpha = 0.01
The Train log loss is: 0.5098581958138237
The Test log loss is: 0.521919486355566
=====
values of alpha = 0.1
The Train log loss is: 0.5314422133355698
The Test log loss is: 0.540589236217171
=====
values of alpha = 1
The Train log loss is: 0.5639767546292684
The Test log loss is: 0.5695336195940539
=====
values of alpha = 10
The Train log loss is: 0.5840678011399411
The Test log loss is: 0.5902832959597724
```



In [224]:

```
fig, ax = plt.subplots()
ax.plot(alpha, log_error_test, label="test")
# for i, txt in enumerate(np.round(log_error_array,3)):
#     ax.annotate((alpha[i],np.round(txt,3)), (alpha[i],log_error_array[i]))

ax.plot(alpha, log_error_tr, label="train")
# for i, txt in enumerate(np.round(log_error_array,3)):
#     ax.annotate((alpha[i],np.round(txt,3)), (alpha[i],log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.legend()
plt.show()
```



In [225]:

```
best_alpha = np.argmin(log_error_test)
alpha[best_alpha]
```

Out[225]:

0.001

In [226]:

```
clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l1', loss='hinge', random_state=42)
clf.fit(x_tr_set2, y_tr_set2)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(x_tr_set2, y_tr_set2)

predict_y = sig_clf.predict_proba(x_tr_set2)
loss_tr = log_loss(y_tr_set2, predict_y, labels=clf.classes_, eps=1e-15)
print('For values of best alpha = ', alpha[best_alpha], "The train log loss is:", log_loss(y_tr_set2, predict_y, labels=clf.classes_, eps=1e-15))

predict_y = sig_clf.predict_proba(x_test_set2)
loss_test = log_loss(y_test_set2, predict_y, labels=clf.classes_, eps=1e-15)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:", log_loss(y_test_set2, predict_y, labels=clf.classes_, eps=1e-15))


predicted_y = np.argmax(predict_y,axis=1)
print("Total number of data points :", len(predicted_y))


plot_confusion_matrix(y_test_set2, predicted_y)
```

For values of best alpha = 0.001 The train log loss is: 0.5389807644172705

For values of best alpha = 0.001 The test log loss is: 0.551159201065387

Total number of data points : 3000





In [227]:

```
# summary.field_names = ["Vectorizer", "Model", "Best Hyperparameters", "Log loss on Train", "Log
loss on Test"]
summary.add_row(['TFIDF', 'Linear SVM by iterative search', "L2 Norm & alpha = " +
str(alpha[best_alpha]) , loss_tr, loss_test ])
print(summary)
```

Vectorizer	Model	Best Hyperparameters	Log loss on Train	Log loss on Test
TFIDF_w2v	Random dumb model	-	-	0.8909857413369713
TFIDF	Random dumb model	-	-	0.8839580037167362
TFIDF_w2v	Logistic Regression by iterative search	alpha = 0.1	-	0.9077642168454928
TFIDF	Logistic Regression by iterative search	alpha = 0.001	-	0.5181022850363174
TFIDF_w2v	Linear SVM by iterative search	L1 Norm & alpha = 0.0001	-	0.5004491052297945
TFIDF_w2v	Linear SVM by iterative search	L2 Norm & alpha = 0.1	-	0.5149092374073378
TFIDF	Linear SVM by iterative search	L1 Norm & alpha = 1e-05	-	0.5528737395236669
TFIDF	Linear SVM by iterative search	L2 Norm & alpha = 0.001	-	0.49838297495456246
TFIDF	Linear SVM by iterative search	L2 Norm & alpha = 0.001	-	0.5148041636287534
TFIDF	Linear SVM by iterative search	L2 Norm & alpha = 0.001	-	0.5389807644172705
TFIDF	Linear SVM by iterative search	L2 Norm & alpha = 0.001	-	0.551159201065387

## 9.4 XGBoost

### 9.4.1 on set1 TFIDF-W2V

In [228]:

```
# https://www.kaggle.com/phunter/xgboost-with-gridsearchcv
xgb_clf = xgb.XGBClassifier()

parameters = {'objective' : ['binary:logistic'], 'eval_metric' : ['logloss'],
              'max_depth': [1,2,3,4],
              'learning_rate': [0.001, 0.01, 0.1, 0.2, 0.3],
              }

d_train = xgb.DMatrix(x_tr_set1, label=y_tr_set1)
d_test = xgb.DMatrix(x_test_set1, label=y_test_set1)

watchlist = [(d_train, 'train'), (d_test, 'valid')]
```

In [229]:

```
clf = GridSearchCV(xgb_clf, parameters, n_jobs = -1,
                  cv=3,
                  scoring='neg_log_loss',
                  verbose=2, refit=True)
```

In [230]:

```
clf.fit(x_tr_set1, y_tr_set1)
```

Fitting 3 folds for each of 24 candidates, totalling 72 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 2 concurrent workers.  
[Parallel(n_jobs=-1)]: Done 37 tasks      | elapsed: 2.5min  
[Parallel(n_jobs=-1)]: Done 72 out of 72 | elapsed: 4.4min finished
```

Out[230]:

```
GridSearchCV(cv=3, error_score=nan,  
             estimator=XGBClassifier(base_score=0.5, booster='gbtree',  
                                     colsample_bylevel=1, colsample_bynode=1,  
                                     colsample_bytree=1, gamma=0,  
                                     learning_rate=0.1, max_delta_step=0,  
                                     max_depth=3, min_child_weight=1,  
                                     missing=None, n_estimators=100, n_jobs=1,  
                                     nthread=None, objective='binary:logistic',  
                                     random_state=0, reg_alpha=0, reg_lambda=1,  
                                     scale_pos_weight=1, seed=None, silent=None,  
                                     subsample=1, verbosity=1),  
             iid='deprecated', n_jobs=-1,  
             param_grid={'eval_metric': ['logloss'],  
                         'learning_rate': [0.001, 0.01, 0.1, 0.2, 0, 3],  
                         'max_depth': [1, 2, 3, 4],  
                         'objective': ['binary:logistic']},  
             pre_dispatch='2*n_jobs', refit=True, return_train_score=False,  
             scoring='neg_log_loss', verbose=2)
```

In [231]:

```
for param_name in sorted(clf.best_params_.keys()):  
    print((param_name, clf.best_params_[param_name]))
```

```
('eval_metric', 'logloss')  
('learning_rate', 0.1)  
('max_depth', 4)  
('objective', 'binary:logistic')
```

In [232]:

```
best_lr = clf.best_params_['learning_rate']  
best_max_depth = clf.best_params_['max_depth']
```

In [233]:

```
import xgboost as xgb  
params = {}  
params['objective'] = 'binary:logistic'  
params['eval_metric'] = 'logloss'  
params['learning_rate'] = best_lr  
params['max_depth'] = best_max_depth  
  
d_train = xgb.DMatrix(x_tr_set1, label=y_tr_set1)  
d_test = xgb.DMatrix(x_test_set1, label=y_test_set1)  
  
watchlist = [(d_train, 'train'), (d_test, 'valid')]  
  
bst = xgb.train(params, d_train, 400, watchlist, early_stopping_rounds=20, verbose_eval=10)  
  
xgdmatrix = xgb.DMatrix(x_tr_set1, y_tr_set1)  
predict_y = bst.predict(d_test)
```

```
[0] train-logloss:0.66132 valid-logloss:0.66267  
Multiple eval metrics have been passed: 'valid-logloss' will be used for early stopping.
```

Will train until valid-logloss hasn't improved in 20 rounds.

```
[10] train-logloss:0.512494 valid-logloss:0.522546  
[20] train-logloss:0.458697 valid-logloss:0.479964  
[30] train-logloss:0.431779 valid-logloss:0.462746  
[40] train-logloss:0.414161 valid-logloss:0.455426  
[50] train-logloss:0.398559 valid-logloss:0.45054  
[60] train-logloss:0.388887 valid-logloss:0.446608
```

```
[60] train-logloss:0.382891 valid-logloss:0.446608
[70] train-logloss:0.366914 valid-logloss:0.443654
[80] train-logloss:0.355275 valid-logloss:0.441785
[90] train-logloss:0.344767 valid-logloss:0.440792
[100] train-logloss:0.332641 valid-logloss:0.440071
[110] train-logloss:0.321628 valid-logloss:0.439307
[120] train-logloss:0.311207 valid-logloss:0.437578
[130] train-logloss:0.298065 valid-logloss:0.436514
[140] train-logloss:0.289942 valid-logloss:0.436009
[150] train-logloss:0.280797 valid-logloss:0.435821
[160] train-logloss:0.272584 valid-logloss:0.435975
[170] train-logloss:0.264299 valid-logloss:0.434842
[180] train-logloss:0.255647 valid-logloss:0.434292
[190] train-logloss:0.248177 valid-logloss:0.434594
[200] train-logloss:0.241086 valid-logloss:0.434431
Stopping. Best iteration:
[183] train-logloss:0.252949 valid-logloss:0.43408
```

In [234]:


```
print("The test log loss is:", log_loss(y_test_set1, predict_y, eps=1e-15))
```

The test log loss is: 0.43423951045549375

In [235]:

```
predicted_y = np.array(predict_y>0.5, dtype=int)
print("Total number of data points :", len(predicted_y))
plot_confusion_matrix(y_test_set1, predicted_y)
```

Total number of data points : 3000



In [236]:

```
predict_y = bst.predict(d_train)
loss_tr = log_loss(y_tr_set1, predict_y, eps=1e-15)
predict_y = bst.predict(d_test)
loss_test = log_loss(y_test_set1, predict_y, eps=1e-15)
```

In [237]:

```
summary.add_row(['TFIDF_w2v', 'XGBoost by GridSearchCV', 'learning_rate = '+str(best_lr)+' &
max_depth = '+str(best_max_depth) , loss_tr, loss_test ])
print(summary)
```

Vectorizer		Model	Best Hyperparameters	Log
loss on Train	Log loss on Test			
TFIDF_w2v	Random dumb model	-	0.8	
9857413369713	0.8839580037167362			
TFIDF	Random dumb model	-	0.8	
1386222435589	0.9077642168454928			
TFIDF_w2v	Logistic Regression by iterative search	alpha = 0.1		
0.5181022850363174	0.5504720745790216			
TFIDF	Logistic Regression by iterative search	alpha = 0.001	0.8	

TFIDF	Logistic Regression by iterative search	alpha = 0.001	0.5
04491052297945	0.5149092374073378		
TFIDF_w2v	Linear SVM by iterative search	L1 Norm & alpha = 0.0001	0.5
27068330818554	0.5528737395236669		
TFIDF_w2v	Linear SVM by iterative search	L2 Norm & alpha = 0.1	0.5
956886197710739	0.6080006713561955		
TFIDF	Linear SVM by iterative search	L1 Norm & alpha = 1e-05	
0.49838297495456246	0.5148041636287534		
TFIDF	Linear SVM by iterative search	L2 Norm & alpha = 0.001	
0.5389807644172705	0.551159201065387		
TFIDF_w2v	XGBoost by GridSearchCV	learning_rate = 0.1 & max_depth = 4	0.23
898285190906504	0.43423951045549375		

## 9.4.2 on set2 TFIDF

In [238]:

```
# https://www.kaggle.com/phunter/xgboost-with-gridsearchcv
xgb_clf = xgb.XGBClassifier()

parameters = {'objective' : ['binary:logistic'], 'eval_metric' : ['logloss'],
              'max_depth': [1,2,3,4],
              'learning_rate': [0.001, 0.01, 0.1, 0.2, 0.3],
              }
d_train = xgb.DMatrix(x_tr_set2, label=y_tr_set2)
d_test = xgb.DMatrix(x_test_set2, label=y_test_set2)

watchlist = [(d_train, 'train'), (d_test, 'valid')]
```

In [239]:

```
clf = GridSearchCV(xgb_clf, parameters, n_jobs = -1,
                  cv=3,
                  scoring='neg_log_loss',
                  verbose=2, refit=True)
```

In [240]:

```
clf.fit(x_tr_set2, y_tr_set2)
```

Fitting 3 folds for each of 24 candidates, totalling 72 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 2 concurrent workers.
[Parallel(n_jobs=-1)]: Done 37 tasks | elapsed: 1.0min
[Parallel(n_jobs=-1)]: Done 72 out of 72 | elapsed: 1.8min finished
```

Out[240]:

```
GridSearchCV(cv=3, error_score=nan,
             estimator=XGBClassifier(base_score=0.5, booster='gbtree',
                                     colsample_bylevel=1, colsample_bynode=1,
                                     colsample_bytree=1, gamma=0,
                                     learning_rate=0.1, max_delta_step=0,
                                     max_depth=3, min_child_weight=1,
                                     missing=None, n_estimators=100, n_jobs=1,
                                     nthread=None, objective='binary:logistic',
                                     random_state=0, reg_alpha=0, reg_lambda=1,
                                     scale_pos_weight=1, seed=None, silent=None,
                                     subsample=1, verbosity=1),
             iid='deprecated', n_jobs=-1,
             param_grid={'eval_metric': ['logloss'],
                         'learning_rate': [0.001, 0.01, 0.1, 0.2, 0.3],
                         'max_depth': [1, 2, 3, 4],
                         'objective': ['binary:logistic']},
             pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
             scoring='neg_log_loss', verbose=2)
```

In [241]:

```
for param_name in sorted(clf.best_params_.keys()):
    print((param_name, clf.best_params_[param_name]))
```

```
('eval_metric', 'logloss')
('learning_rate', 0.2)
('max_depth', 4)
('objective', 'binary:logistic')
```

In [242]:

```
best_lr = clf.best_params_['learning_rate']
best_max_depth = clf.best_params_['max_depth']
```

In [243]:

```
import xgboost as xgb
params = {}
params['objective'] = 'binary:logistic'
params['eval_metric'] = 'logloss'
params['learning_rate'] = best_lr
params['max_depth'] = best_max_depth

d_train = xgb.DMatrix(x_tr_set2, label=y_tr_set2)
d_test = xgb.DMatrix(x_test_set2, label=y_test_set2)

watchlist = [(d_train, 'train'), (d_test, 'valid')]

bst = xgb.train(params, d_train, 400, watchlist, early_stopping_rounds=20, verbose_eval=10)

xgdmatrix = xgb.DMatrix(x_tr_set2, y_tr_set2)
predict_y = bst.predict(d_test)
```

```
[0] train-logloss:0.633313 valid-logloss:0.635351
Multiple eval metrics have been passed: 'valid-logloss' will be used for early stopping.
```

Will train until valid-logloss hasn't improved in 20 rounds.

```
[10] train-logloss:0.458116 valid-logloss:0.477098
[20] train-logloss:0.423672 valid-logloss:0.453158
[30] train-logloss:0.40006 valid-logloss:0.441863
[40] train-logloss:0.381075 valid-logloss:0.434494
[50] train-logloss:0.362261 valid-logloss:0.429763
[60] train-logloss:0.347543 valid-logloss:0.426033
[70] train-logloss:0.335782 valid-logloss:0.422541
[80] train-logloss:0.323992 valid-logloss:0.421173
[90] train-logloss:0.316225 valid-logloss:0.421302
[100] train-logloss:0.307727 valid-logloss:0.420663
[110] train-logloss:0.300978 valid-logloss:0.420736
[120] train-logloss:0.292785 valid-logloss:0.41968
[130] train-logloss:0.285708 valid-logloss:0.419546
[140] train-logloss:0.278457 valid-logloss:0.421014
[150] train-logloss:0.272805 valid-logloss:0.421247
Stopping. Best iteration:
[133] train-logloss:0.282118 valid-logloss:0.418883
```

In [244]:


```
print("The test log loss is:", log_loss(y_test_set2, predict_y, eps=1e-15))
```

The test log loss is: 0.42154032061294855

In [245]:

```
predicted_y = np.array(predict_y > 0.5, dtype=int)
print("Total number of data points :", len(predicted_y))
plot_confusion_matrix(y_test_set2, predicted_y)
```

Total number of data points : 3000



In [246]:

```
predict_y = bst.predict(d_train)
loss_tr = log_loss(y_tr_set2, predict_y, eps=1e-15)
predict_y = bst.predict(d_test)
loss_test = log_loss(y_test_set2, predict_y, eps=1e-15)
```

In [247]:

```
summary.add_row(['TFIDF', 'XGBoost by GridSearchCV', 'learning_rate = '+str(best_lr)+' & max_depth = '+str(best_max_depth) , loss_tr, loss_test ])
print(summary)
```

Vectorizer	Model	Best Hyperparameters	Log
loss on Train	Log loss on Test		
TFIDF_w2v	Random dumb model	-	0.8
9857413369713	0.8839580037167362		
TFIDF	Random dumb model	-	0.8
1386222435589	0.9077642168454928		
TFIDF_w2v	Logistic Regression by iterative search	alpha = 0.1	
0.5181022850363174	0.5504720745790216		
TFIDF	Logistic Regression by iterative search	alpha = 0.001	0.5
04491052297945	0.5149092374073378		
TFIDF_w2v	Linear SVM by iterative search	L1 Norm & alpha = 0.0001	0.5
27068330818554	0.5528737395236669		
TFIDF_w2v	Linear SVM by iterative search	L2 Norm & alpha = 0.1	0.5
956886197710739	0.6080006713561955		
TFIDF	Linear SVM by iterative search	L1 Norm & alpha = 1e-05	
0.49838297495456246	0.5148041636287534		
TFIDF	Linear SVM by iterative search	L2 Norm & alpha = 0.001	
0.5389807644172705	0.551159201065387		
TFIDF_w2v	XGBoost by GridSearchCV	learning_rate = 0.1 & max_depth = 4	0.23
898285190906504	0.43423951045549375		
TFIDF	XGBoost by GridSearchCV	learning_rate = 0.2 & max_depth = 4	0.27
033581958191283	0.42154032061294855		

## 10.Conclusion

In [248]:

```
print(summary)
```

Vectorizer	Model	Best Hyperparameters	Log
loss on Train	Log loss on Test		
TFIDF_w2v	Random dumb model	-	0.8
9857413369713	0.8839580037167362		
TFIDF	Random dumb model	-	0.8
1386222435589	0.9077642168454928		
TFIDF_w2v	Logistic Regression by iterative search	alpha = 0.1	
0.5181022850363174	0.5504720745790216		

◀ ▶

## 11.Observations and in detail Summary

1. From EDA of the data we observed that:
  - The dataset is imbalanced with 63-37% approx to non duplicate to duplicate question pairs
  - There is very less percentage of Unique question being repeated.
  - There are no repeated questions pairs in Dataset.
  - There are no NULL entries.
2. Basic Feature Engineering
  - We created 11 Features before preprocessing the data and observed the following:
    - There are quite a lot of questions with high word similarity
    - The average word share and Common no. of words of qid1 and qid2 is more when they are duplicate(Similar)
    - The distributions of the word\_Common feature in similar and non-similar questions are highly overlapping
3. Preprocessing of the data
  - Removing html tags
  - Removing Punctuations
  - Performing stemming
  - Removing Stopwords
  - Expanding contractions etc.
4. Advance Feature Engineering
  - We created 15 Advance Features after preprocessing the data and observed the following:
    - There are few words which occurs frequently in both the cases when the questions are similar and not similar.
    - From the pair plots of few features, It can be observed that there are few non duplicate pairs of questions which are linearly spearable, although there is a lot of overlapping of distribution as well.
    - The TSNE plot very well proves the above point that there are few clusters of non duplicate pairs which are not overlapping at all with clusters of duplicate pairs.
    - The TSNE plot also shows that there are majority of clusters of non duplicate pairs which are overlapping with clusters of duplicate pairs.
5. Splitting of Data
  - Data is splitted into 70-30 ratio into Train and Test
6. Featurizing Text Data
  - Created TFIDF for Train and Test data by fitting vectorizer on train data and transforming it on train and test, also limiting the max\_fetures to 100.
  - created TFIDF weighted Word to Vec for Train and Test data appropriately.
7. Creating Data ready for Training Models
  - Set1: comprises of all 11 basic features + 15 advance features + 96 features of TFIDF weighted word to vec of Que 1 + 96 features of TFIDF weighted word to vec of Que 2.
  - Set2: comprises of all 11 basic features + 15 advance features + 100 features of TFIDF of Que 1 + 100 features of TFIDF of Que 2.
8. Training Models
  - Random Dumb Model gives log loss of 0.89 and 0.90 on Set1 and Set2 respectively.
  - Logistic Regression works better on Set2 as compared to Set1 Data.
  - Linear SVM works better on Set2 as compared to Set1 Data.
  - XGBoost works marginally better on Set2 as compared to Set1 Data.